

# knn

October 9, 2022

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'ENPM809K Assignments/assignment1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/datasets
/content/drive/My Drive/ENPM809K Assignments/assignment1
```

## 1 k-Nearest Neighbor (kNN) exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

The kNN classifier consists of two stages:

- During training, the classifier takes the training data and simply remembers it
- During testing, kNN classifies every test image by comparing to all training images and transferring the labels of the k most similar training examples
- The value of k is cross-validated

In this exercise you will implement these steps and understand the basic Image Classification pipeline, cross-validation, and gain proficiency in writing efficient, vectorized code.

```
[ ]: # Run some setup code for this notebook.

import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
↪notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
↪autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

[ ]: # Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

# Cleaning up variables to prevent loading data multiple times (which may cause
↪memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
# PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
# PLEASE DO NOT MODIFY THE MARKERS
print('~~~~~')
```

```

|||||
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
.....

```

```

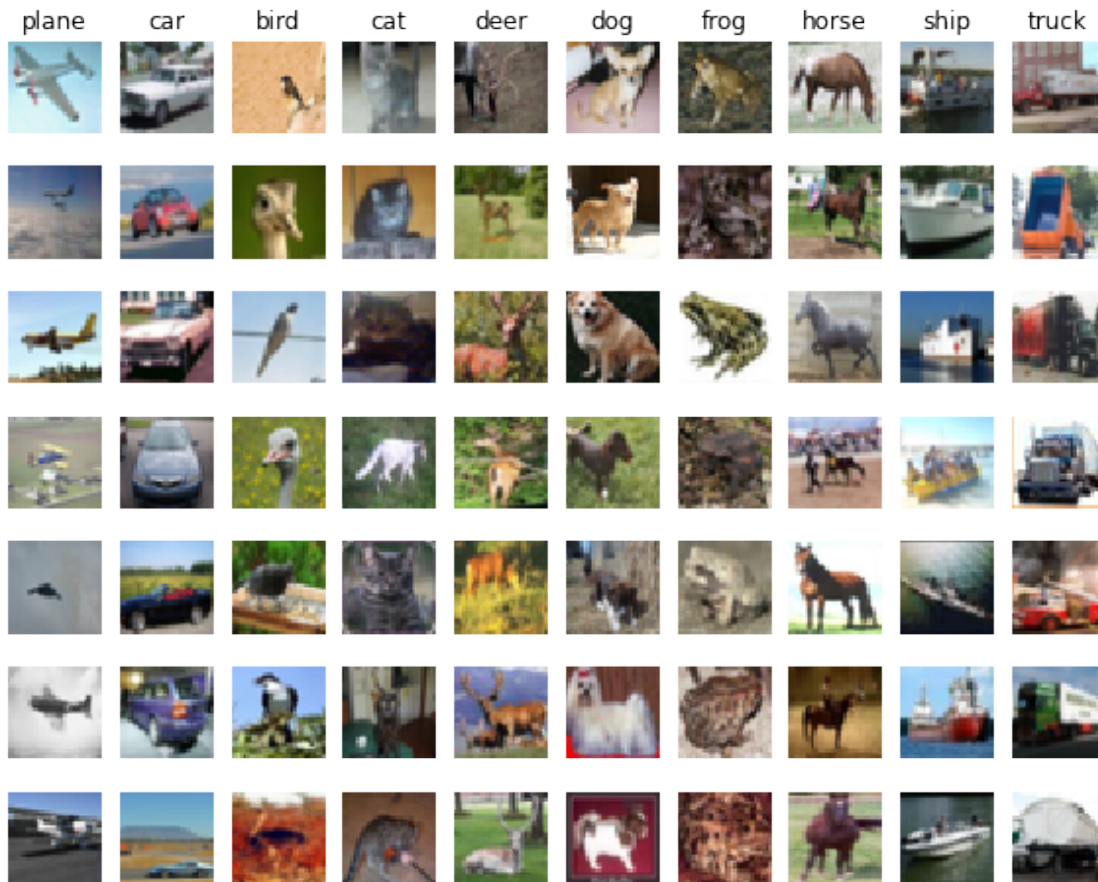
[ ]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
# PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
↪'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
print('.....')

```

```

|||||

```



```
[ ]: # Subsample the data for more efficient code execution in this exercise
num_training = 5000
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]

num_test = 500
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]

# Reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
# PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
print(X_train.shape, X_test.shape)
print('-----')
```

```

.....
|||||
(5000, 3072) (500, 3072)
.....

```

```

[ ]: from cs231n.classifiers import KNearestNeighbor
# PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
# Create a kNN classifier instance.
# Remember that training a kNN classifier is a noop:
# the Classifier simply remembers the data and does no further processing
classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
print('.....')

```

```

|||||
.....

```

We would now like to classify the test data with the kNN classifier. Recall that we can break down this process into two steps:

1. First we must compute the distances between all test examples and all train examples.
2. Given these distances, for each test example we find the k nearest examples and have them vote for the label

Lets begin with computing the distance matrix between all training and test examples. For example, if there are **Ntr** training examples and **Nte** test examples, this stage should result in a **Nte x Ntr** matrix where each element (i,j) is the distance between the i-th test and j-th train example.

**Note:** For the three distance computations that we require you to implement in this notebook, you may not use the `np.linalg.norm()` function that numpy provides.

First, open `cs231n/classifiers/k_nearest_neighbor.py` and implement the function `compute_distances_two_loops` that uses a (very inefficient) double loop over all pairs of (test, train) examples and computes the distance matrix one element at a time.

```

[ ]: # Open cs231n/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.
# PLEASE DO NOT MODIFY THE MARKERS
print('|||||')

# Test your implementation:
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')


```

```

|||||
(500, 5000)
.....

```

```
[ ]: # We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
# PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
plt.imshow(dists, interpolation='none')
plt.show()
# PLEASE DO NOT MODIFY THE MARKERS
print('-----')
```



```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

### Inline Question 1

Notice the structured patterns in the distance matrix, where some rows or columns are visible brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

- What in the data is the cause behind the distinctly bright rows?
- What causes the columns?

*Your Answer :* The distinctly bright rows could signify that test images in that row are distinctly different from the corresponding training images. Since black denotes lower distance between images, meaning they could be of the same class. Now since its all one single colour, black, we can't say which image belongs to which class and so on. So we cannot infer much from either the black rows and columns.

```
[ ]: # Now implement the function predict_labels and run the code below:
# We use k = 1 (which is Nearest Neighbor).
y_test_pred = classifier.predict_labels(dists, k=1)

# Compute and print the fraction of correctly predicted examples
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
# PLEASE DO NOT MODIFY THE MARKERS
```

```

print('||||||||||||||||||||||||||||||||||||||||')
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')

```

```

||||||||||||||||||||||||||||||||||||||||
Got 137 / 500 correct => accuracy: 0.274000
.....

```

You should expect to see approximately 27% accuracy. Now lets try out a larger k, say k = 5:

```

[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('||||||||||||||||||||||||||||||||||||||||')
y_test_pred = classifier.predict_labels(dists, k=5)
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
# PLEASE DO NOT MODIFY THE MARKERS
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
print('.....')

```

```

||||||||||||||||||||||||||||||||||||||||
Got 139 / 500 correct => accuracy: 0.278000
.....

```

You should expect to see a slightly better performance than with k = 1.

## Inline Question 2

We can also use other distance metrics such as L1 distance. For pixel values  $p_{ij}^{(k)}$  at location  $(i, j)$  of some image  $I_k$ ,

the mean  $\mu$  across all pixels over all images is

$$\mu = \frac{1}{nhw} \sum_{k=1}^n \sum_{i=1}^h \sum_{j=1}^w p_{ij}^{(k)}$$

And the pixel-wise mean  $\mu_{ij}$  across all images is

$$\mu_{ij} = \frac{1}{n} \sum_{k=1}^n p_{ij}^{(k)}.$$

The general standard deviation  $\sigma$  and pixel-wise standard deviation  $\sigma_{ij}$  is defined similarly.

Which of the following preprocessing steps will not change the performance of a Nearest Neighbor classifier that uses L1 distance? Select all that apply. 1. Subtracting the mean  $\mu$  ( $\tilde{p}_{ij}^{(k)} = p_{ij}^{(k)} - \mu$ .) 2. Subtracting the per pixel mean  $\mu_{ij}$  ( $\tilde{p}_{ij}^{(k)} = p_{ij}^{(k)} - \mu_{ij}$ .) 3. Subtracting the mean  $\mu$  and dividing by the standard deviation  $\sigma$ . 4. Subtracting the pixel-wise mean  $\mu_{ij}$  and dividing by the pixel-wise standard deviation  $\sigma_{ij}$ . 5. Rotating the coordinate axes of the data.

*Your Answer* : 1, 3

### *Your Explanation :*

1 -> For this, since in L1 distance, since we are taking absolute differences between two images and summing them all up, taking their means and subtracting from each image will lead to the means to cancel each other out since they will be equal.

3 -> The subtracting the mean has the similar explanation. In addition to it, by dividing by the standard deviation, we will only have a scaling of the distance by a factor which will not result in varied performance of the algorithm.

```
[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('||||||||||||||||||||||||||||||||||||||||')
# Now lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
dists_one = classifier.compute_distances_one_loop(X_test)

# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words,
# →reshape
# the matrices into vectors and compute the Euclidean distance between them.
difference = np.linalg.norm(dists - dists_one, ord='fro')
print('One loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')
```

```
||||||||||||||||||||||||||||||||||||||||
One loop difference was: 0.000000
Good! The distance matrices are the same
.....
```

```
[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('||||||||||||||||||||||||||||||||||||||||')
# Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)

# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('No loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
```



```

    print('Uh-oh! The distance matrices are different')
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')

```

```

|||||
No loop difference was: 0.000000
Good! The distance matrices are the same
.....

```

```

[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
# Let's compare how fast the implementations are
def time_function(f, *args):
    """
    Call a function f with args and return the time (in seconds) that it took
    ↪to execute.
    """
    import time
    tic = time.time()
    f(*args)
    toc = time.time()
    return toc - tic

two_loop_time = time_function(classifier.compute_distances_two_loops, X_test)
print('Two loop version took %f seconds' % two_loop_time)

one_loop_time = time_function(classifier.compute_distances_one_loop, X_test)
print('One loop version took %f seconds' % one_loop_time)

no_loop_time = time_function(classifier.compute_distances_no_loops, X_test)
print('No loop version took %f seconds' % no_loop_time)

# You should see significantly faster performance with the fully vectorized
↪implementation!

# NOTE: depending on what machine you're using,
# you might not see a speedup when you go from two loops to one loop,
# and might even see a slow-down.
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')

```

```

|||||
Two loop version took 56.621530 seconds
One loop version took 55.121733 seconds
No loop version took 0.578792 seconds
.....

```

### 1.0.1 Cross-validation

We have implemented the k-Nearest Neighbor classifier but we set the value  $k = 5$  arbitrarily. We will now determine the best value of this hyperparameter with cross-validation.

```
[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
num_folds = 5
k_choices = [1, 3, 5, 8, 10, 12, 15, 20, 50, 100]

X_train_folds = []
y_train_folds = []
#####
# TODO:
# Split up the training data into folds. After splitting, X_train_folds and
# y_train_folds should each be lists of length num_folds, where
# y_train_folds[i] is the label vector for the points in X_train_folds[i].
# Hint: Look up the numpy array_split function.
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

#First, split the training images into the corresponding number of folds
X_train_folds = np.array_split(X_train, num_folds)
#Second, split the training labels into the corresponding number of folds
y_train_folds = np.array_split(y_train, num_folds)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# A dictionary holding the accuracies for different values of k that we find
# when running cross-validation. After running cross-validation,
# k_to_accuracies[k] should be a list of length num_folds giving the different
# accuracy values that we found when using that value of k.
k_to_accuracies = {}

#####
# TODO:
# Perform k-fold cross validation to find the best value of k. For each
# possible value of k, run the k-nearest-neighbor algorithm num_folds times,
# where in each case you use all but one of the folds as training data and the
# last fold as a validation set. Store the accuracies for all fold and all
# values of k in the k_to_accuracies dictionary.
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

#First, we iterate through each of the presented k_choices
```

```

#Second, initiate an empty list giving different accuracies for each k value
for k in k_choices:
    k_to_accuracies[k] = []
    #Third, we choose the present fold for cross validation by traversing through
    #the number of folds
    for i in range(num_folds):
        #Fourth, we prepare the training images
        #All folds except the current fold should be treated as training set
        current_X_train_fold = [] #initiate an empty list to store all the data
        for idxss, each_x in enumerate(X_train_folds): #traverse through it
            if i != idxss: #choose only those folds which aren't the same as the
→current fold
                current_X_train_fold.append(each_x) #append them all together and
→concatenate them to get the current training fold
        current_X_train_fold = np.concatenate(current_X_train_fold)

        #Now, we prepare the labels for the train image in the same manner
        current_y_train_fold = []
        for idyss, each_y in enumerate(y_train_folds):
            if i != idyss:
                current_y_train_fold.append(each_y)
        current_y_train_fold = np.concatenate(current_y_train_fold)

        #Fifth, we train the algorithm for k-nn classification
        classifier.train(current_X_train_fold, current_y_train_fold)

        #Sixth, we use this to predict our answers for the current fold
        predicted_y_current_fold = classifier.predict(X_train_folds[i], k, 0)

        #Seventh, we now check for accuracies of our prediction
        prediction_accuracy = np.mean(predicted_y_current_fold == y_train_folds[i])

        #Last, we just add the accuracy as a value to that particular key in the
→accuracies dictionary
        k_to_accuracies[k].append(prediction_accuracy)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out the computed accuracies
for k in sorted(k_to_accuracies):
    for accuracy in k_to_accuracies[k]:
        print('k = %d, accuracy = %f' % (k, accuracy))
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')

```

||||||||||||||||||||||||||||||||||||||||||||||||||||||||

k = 1, accuracy = 0.263000  
k = 1, accuracy = 0.257000  
k = 1, accuracy = 0.264000  
k = 1, accuracy = 0.278000  
k = 1, accuracy = 0.266000  
k = 3, accuracy = 0.239000  
k = 3, accuracy = 0.249000  
k = 3, accuracy = 0.240000  
k = 3, accuracy = 0.266000  
k = 3, accuracy = 0.254000  
k = 5, accuracy = 0.248000  
k = 5, accuracy = 0.266000  
k = 5, accuracy = 0.280000  
k = 5, accuracy = 0.292000  
k = 5, accuracy = 0.280000  
k = 8, accuracy = 0.262000  
k = 8, accuracy = 0.282000  
k = 8, accuracy = 0.273000  
k = 8, accuracy = 0.290000  
k = 8, accuracy = 0.273000  
k = 10, accuracy = 0.265000  
k = 10, accuracy = 0.296000  
k = 10, accuracy = 0.276000  
k = 10, accuracy = 0.284000  
k = 10, accuracy = 0.280000  
k = 12, accuracy = 0.260000  
k = 12, accuracy = 0.295000  
k = 12, accuracy = 0.279000  
k = 12, accuracy = 0.283000  
k = 12, accuracy = 0.280000  
k = 15, accuracy = 0.252000  
k = 15, accuracy = 0.289000  
k = 15, accuracy = 0.278000  
k = 15, accuracy = 0.282000  
k = 15, accuracy = 0.274000  
k = 20, accuracy = 0.270000  
k = 20, accuracy = 0.279000  
k = 20, accuracy = 0.279000  
k = 20, accuracy = 0.282000  
k = 20, accuracy = 0.285000  
k = 50, accuracy = 0.271000  
k = 50, accuracy = 0.288000  
k = 50, accuracy = 0.278000  
k = 50, accuracy = 0.269000  
k = 50, accuracy = 0.266000  
k = 100, accuracy = 0.256000  
k = 100, accuracy = 0.270000

```

k = 100, accuracy = 0.263000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.263000
.....

```

```

[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
# plot the raw observations
for k in k_choices:
    accuracies = k_to_accuracies[k]
    plt.scatter([k] * len(accuracies), accuracies)

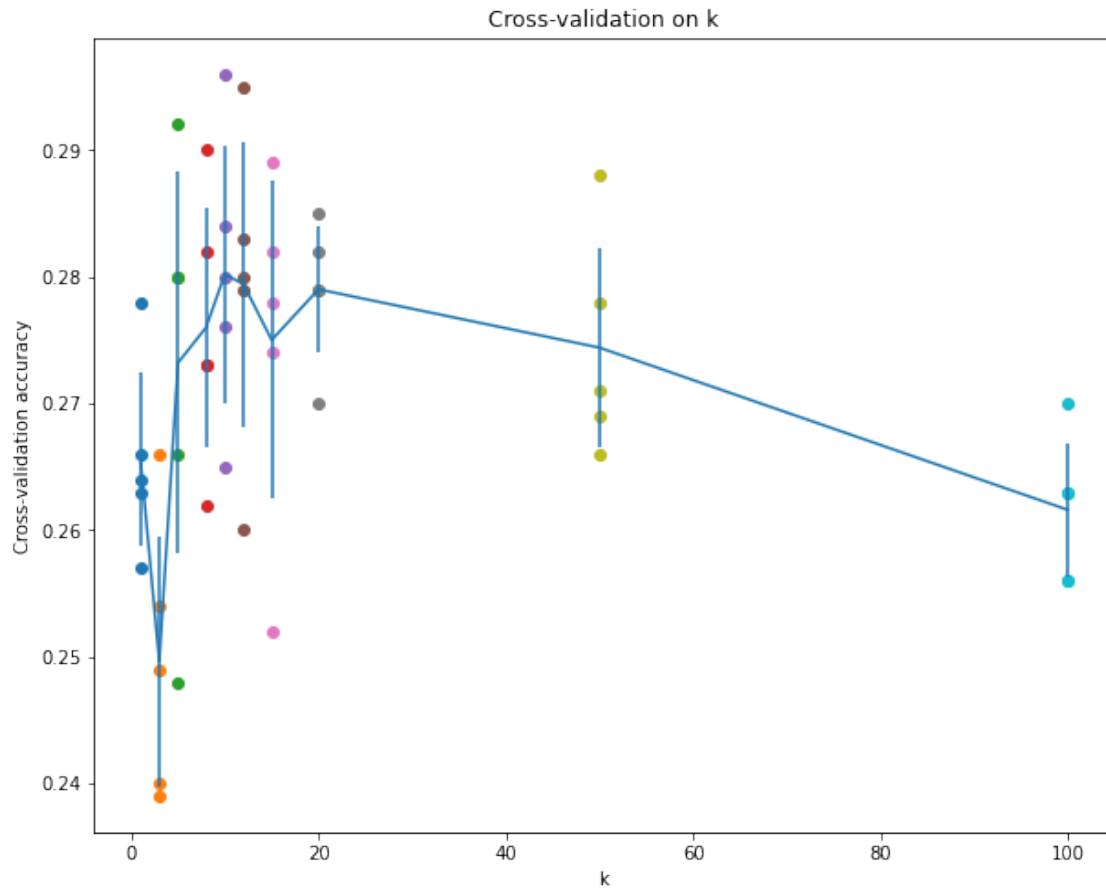
# plot the trend line with error bars that correspond to standard deviation
accuracies_mean = np.array([np.mean(v) for k,v in sorted(k_to_accuracies.
    ↳items())])
accuracies_std = np.array([np.std(v) for k,v in sorted(k_to_accuracies.
    ↳items())])
plt.errorbar(k_choices, accuracies_mean, yerr=accuracies_std)
plt.title('Cross-validation on k')
plt.xlabel('k')
plt.ylabel('Cross-validation accuracy')
plt.show()
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')

```

```

|||||

```



.....

```
[ ]: # PLEASE DO NOT MODIFY THE MARKERS
print('|||||')
# Based on the cross-validation results above, choose the best value for k,
# retrain the classifier using all the training data, and test it on the test
# data. You should be able to get above 28% accuracy on the test data.
best_k = 1

#Here, we need to choose the k with the best accuracy
#First, choose index of k which has the highest accuracy
#Second, from the list of k_choices, choose k with that index
best_k_index = np.argmax(accuracies_mean)
best_k = k_choices[best_k_index]

classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
y_test_pred = classifier.predict(X_test, k=best_k)
```

```
# Compute and display the accuracy
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
# PLEASE DO NOT MODIFY THE MARKERS
print('.....')
```

```
|||||
Got 141 / 500 correct => accuracy: 0.282000
.....
```

### Inline Question 3

Which of the following statements about  $k$ -Nearest Neighbor ( $k$ -NN) are true in a classification setting, and for all  $k$ ? Select all that apply. 1. The decision boundary of the  $k$ -NN classifier is linear. 2. The training error of a 1-NN will always be lower than or equal to that of 5-NN. 3. The test error of a 1-NN will always be lower than that of a 5-NN. 4. The time needed to classify a test example with the  $k$ -NN classifier grows with the size of the training set. 5. None of the above.

*Your Answer :* 2 & 4

*Your Explanation :* \ 2 -> Because in a Nearest Neighbour (1-NN) classifier, we are only checking with one nearby neighbour, the scope for errors could be lower. Also, if there are no nearest neighbours (incase of outliers), it will be its own nearest neighbour, hence the error will be zero (no chance for any data point to be left unclassified).

4 -> Because in the  $k$ -NN classifier, the training aspect is just about the program remembering the training set once, immaterial of its size. But when it comes to the testing set, since each image in the testing is compared against each image in the training set, the bigger the training set the more computations and thus more time is needed to classify the test example.

# SVM

October 9, 2022

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'ENPM809K Assignments/assignment1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Mounted at /content/drive

/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/datasets

/content/drive/My Drive/ENPM809K Assignments/assignment1

## 1 Multiclass Support Vector Machine exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

In this exercise you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization strength**



- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[ ]: # Run some setup code for this notebook.
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
# notebook rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
  ↳ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

## 1.1 CIFAR-10 Data Loading and Preprocessing

```
[ ]: # Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

# Cleaning up variables to prevent loading data multiple times (which may cause
  ↳ memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```


Training data shape: (50000, 32, 32, 3)

Training labels shape: (50000,)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000,)

```
[ ]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



```
[ ]: # Split the data into train, val, and test sets. In addition we will
# create a small development set as a subset of the training data;
# we can use this for development so our code runs faster.
num_training = 49000
num_validation = 1000
num_test = 1000
num_dev = 500

# Our validation set will be num_validation points from the original
# training set.
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]

# Our training set will be the first num_train points from the original
# training set.
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

# We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# We use the first num_test points of the original test set as our
# test set.
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
```

```
[ ]: # Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('dev data shape: ', X_dev.shape)
```

```
Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

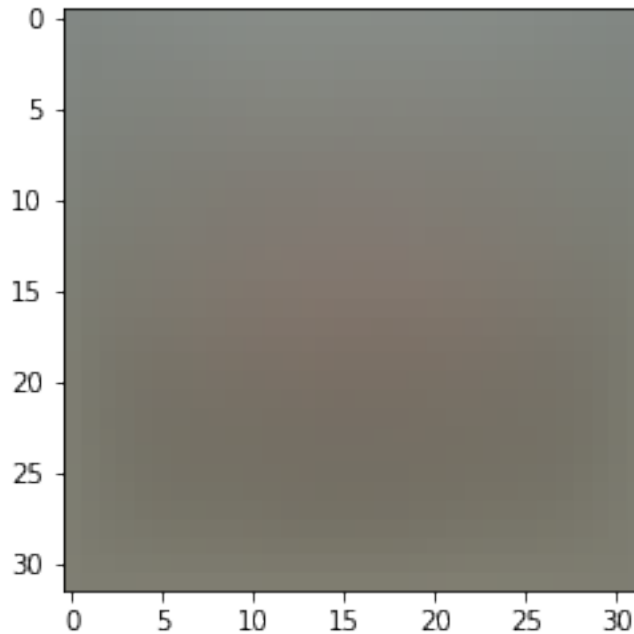
```
[ ]: # Preprocessing: subtract the mean image
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)
print(mean_image[:10]) # print a few of the elements
plt.figure(figsize=(4,4))
plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean_
    ↪ image
plt.show()

# second: subtract the mean image from train and test data
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# third: append the bias dimension of ones (i.e. bias trick) so that our SVM
# only has to worry about optimizing a single weight matrix W.
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)
```

```
[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]
```



(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)

## 1.2 SVM Classifier

Your code for this section will all be written inside `cs231n/classifiers/linear_svm.py`.

As you can see, we have prefilled the function `svm_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

```
[ ]: # Evaluate the naive implementation of the loss we provided for you:
      from cs231n.classifiers.linear_svm import svm_loss_naive
      import time

      # generate a random SVM weight matrix of small numbers
      W = np.random.randn(3073, 10) * 0.0001

      loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      print('loss: %f' % (loss, ))
```

loss: 9.065637

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you

computed. We have provided code that does this for you:

```
[ ]: # Once you've implemented the gradient, recompute it with the code below
# and gradient check it with the function we provided for you

# Compute the loss and its gradient at W.
loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)

# Numerically compute the gradient along several randomly chosen dimensions, and
# compare them with your analytically computed gradient. The numbers should
  ↳ match
# almost exactly along all dimensions.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad)

# do the gradient check once again with regularization turned on
# you didn't forget the regularization gradient did you?
loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad)
```

```
numerical: 5.265230 analytic: 5.265230, relative error: 5.602137e-11
numerical: 32.786009 analytic: 32.786009, relative error: 4.220756e-12
numerical: 10.076683 analytic: 10.076683, relative error: 1.754113e-12
numerical: -13.115585 analytic: -13.115585, relative error: 5.778421e-12
numerical: 10.375401 analytic: 10.408244, relative error: 1.580239e-03
numerical: 6.964807 analytic: 7.003292, relative error: 2.755162e-03
numerical: 13.081231 analytic: 13.081231, relative error: 1.548082e-11
numerical: -9.960189 analytic: -10.004054, relative error: 2.197160e-03
numerical: -18.225744 analytic: -18.225744, relative error: 1.309587e-11
numerical: -20.329009 analytic: -20.329009, relative error: 2.128359e-11
numerical: 28.710120 analytic: 28.710120, relative error: 2.814193e-11
numerical: -5.451096 analytic: -5.451096, relative error: 5.094748e-11
numerical: -33.690499 analytic: -33.650222, relative error: 5.981011e-04
numerical: -26.642634 analytic: -26.655762, relative error: 2.463149e-04
numerical: -38.133790 analytic: -38.099929, relative error: 4.441773e-04
numerical: 9.442944 analytic: 9.442944, relative error: 3.630990e-12
numerical: -14.155677 analytic: -14.155677, relative error: 5.143170e-11
numerical: 27.727049 analytic: 27.727049, relative error: 2.121720e-11
numerical: -0.121644 analytic: -0.121644, relative error: 4.056186e-09
numerical: 14.374267 analytic: 14.429162, relative error: 1.905857e-03
```

### Inline Question 1

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? How would change the margin affect the frequency of this happening? *Hint: the SVM loss function is not strictly speaking differentiable*

*Your Answer*: One probable reason for this could be that the score function at that particular dimension becomes flat, or there maybe a break in the direction of flow (i.e, the function isn't continuous) due to kinks which maybe the results of the max operator used. In such cases, it may not be feasible to differentiate the function at such points. Hence leading to loss in that dimension, i.e dimensionality mismatch. It isn't necessarily a reason of concern because their sub-gradient exists.

```
[ ]: # Next implement the function svm_loss_vectorized; for now only compute the
      ↪ loss;
      # we will implement the gradient in a moment.
      tic = time.time()
      loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))

      from cs231n.classifiers.linear_svm import svm_loss_vectorized
      tic = time.time()
      loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

      # The losses should match but your vectorized implementation should be much
      ↪ faster.
      print('difference: %f' % (loss_naive - loss_vectorized))
```

```
Naive loss: 9.065637e+00 computed in 0.165195s
Vectorized loss: 9.065637e+00 computed in 0.014305s
difference: -0.000000
```

```
[ ]: # Complete the implementation of svm_loss_vectorized, and compute the gradient
      # of the loss function in a vectorized way.

      # The naive implementation and the vectorized implementation should match, but
      # the vectorized version should still be much faster.
      tic = time.time()
      _, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Naive loss and gradient: computed in %fs' % (toc - tic))

      tic = time.time()
      _, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Vectorized loss and gradient: computed in %fs' % (toc - tic))

      # The loss is a single number, so it is easy to compare the values computed
      # by the two implementations. The gradient on the other hand is a matrix, so
      # we use the Frobenius norm to compare them.
      difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
```

```
print('difference: %f' % difference)
```

Naive loss and gradient: computed in 0.166267s  
Vectorized loss and gradient: computed in 0.015663s  
difference: 0.000000

### 1.2.1 Stochastic Gradient Descent


We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss. Your code for this part will be written inside `cs231n/classifiers/linear_classifier.py`.

```
[ ]: # In the file linear_classifier.py, implement SGD in the function
# LinearClassifier.train() and then run it with the code below.
from cs231n.classifiers import LinearSVM
svm = LinearSVM()
tic = time.time()
loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,
                      num_iters=1500, verbose=True)
toc = time.time()
print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 795.830891
iteration 100 / 1500: loss 289.799292
iteration 200 / 1500: loss 109.337541
iteration 300 / 1500: loss 42.545208
iteration 400 / 1500: loss 19.016226
iteration 500 / 1500: loss 9.379888
iteration 600 / 1500: loss 7.044530
iteration 700 / 1500: loss 5.759147
iteration 800 / 1500: loss 5.594958
iteration 900 / 1500: loss 5.072950
iteration 1000 / 1500: loss 5.648617
iteration 1100 / 1500: loss 4.777682
iteration 1200 / 1500: loss 5.176316
iteration 1300 / 1500: loss 5.458756
iteration 1400 / 1500: loss 5.277060
That took 9.313227s
```

```
[ ]: # A useful debugging strategy is to plot the loss as a function of
# iteration number:
plt.plot(loss_hist)
plt.xlabel('Iteration number')
plt.ylabel('Loss value')
plt.show()
```





```
[ ]: # Write the LinearSVM.predict function and evaluate the performance on both the
# training and validation set
y_train_pred = svm.predict(X_train)
print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
y_val_pred = svm.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
```

```
training accuracy: 0.374980
validation accuracy: 0.383000
```

```
[ ]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of about 0.39 on the validation set.

# Note: you may see runtime/overflow warnings during hyper-parameter search.
# This may be caused by extreme values, and is not a bug.

# results is dictionary mapping tuples of the form
# (learning_rate, regularization_strength) to tuples of the form
```

```

# (training_accuracy, validation_accuracy). The accuracy is simply the fraction
# of data points that are correctly classified.
results = {}
best_val = -1    # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation
    ↳rate.

#####
# TODO:
# Write code that chooses the best hyperparameters by tuning on the validation #
# set. For each combination of hyperparameters, train a linear SVM on the #
# training set, compute its accuracy on the training and validation sets, and #
# store these numbers in the results dictionary. In addition, store the best #
# validation accuracy in best_val and the LinearSVM object that achieves this #
# accuracy in best_svm.
#
# Hint: You should use a small value for num_iters as you develop your #
# validation code so that the SVMs don't take much time to train; once you are #
# confident that your validation code works, you should rerun the validation #
# code with a larger value for num_iters.
#####

# Provided as a reference. You may or may not want to change these
    ↳hyperparameters
learning_rates = [1e-7, 5e-5]
regularization_strengths = [2.5e4, 5e4]

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

svm = LinearSVM() #Creating an svm instance to access the classifier

#First, we iterate through each learning rate
#Second, for each learning rate, we iterate through each regularization to find
    ↳the fit
for learn_rate in learning_rates:
    for reg_strength in regularization_strengths:

        #Third, we train our classifier
        loss_function = svm.train(X_train, y_train, learning_rate = learn_rate, reg
    ↳= reg_strength, num_iters=1000)

        #Fourth, we predict the classes on training data and estimate the accuracy
        y_pred_train = svm.predict(X_train)
        training_accuracy = np.mean(y_pred_train == y_train)

        #Fifth, we predict the classes on the validation set and estimate its
    ↳accuracy

```

```

y_pred_validation = svm.predict(X_val)
validation_accuracy = np.mean(y_pred_validation == y_val)

#Next, store these accuracies as dictionaries in the result dict.
results[(learn_rate, reg_strength)] = (training_accuracy,
↪validation_accuracy)

#Lastly, we check if the validation accuracy is better than the best
↪validation accuracy and then update it accordingly
if validation_accuracy > best_val:
    best_val = validation_accuracy
    best_svm = svm #The current SVM will be the best SVM for this iteration

"""Note here, that by increasing/decreasing the number of iteration in our
↪training of the SVM
we can have an direct effect on our validation accuracy"""

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
↪best_val)

```

```

/content/drive/My Drive/ENPM809K
Assignments/assignment1/cs231n/classifiers/linear_svm.py:103: RuntimeWarning:
overflow encountered in double_scalars
    loss += reg * np.sum(W * W)
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:86:
RuntimeWarning: overflow encountered in reduce
    return ufunc.reduce(obj, axis, dtype, out, **passkwargs)
/content/drive/My Drive/ENPM809K
Assignments/assignment1/cs231n/classifiers/linear_svm.py:103: RuntimeWarning:
overflow encountered in multiply
    loss += reg * np.sum(W * W)
/content/drive/My Drive/ENPM809K
Assignments/assignment1/cs231n/classifiers/linear_svm.py:93: RuntimeWarning:
overflow encountered in subtract
    margin = all_scores - correct_class_scores + 1
/content/drive/My Drive/ENPM809K
Assignments/assignment1/cs231n/classifiers/linear_svm.py:93: RuntimeWarning:
invalid value encountered in subtract
    margin = all_scores - correct_class_scores + 1

```

```

/content/drive/My Drive/ENPM809K
Assignments/assignment1/cs231n/classifiers/linear_svm.py:131: RuntimeWarning:
overflow encountered in multiply
    dW += reg * 2 * W # if we diff  $W^2$  we get  $2W$ , apply that to the regularization
term
/content/drive/My Drive/ENPM809K
Assignments/assignment1/cs231n/classifiers/linear_classifier.py:90:
RuntimeWarning: invalid value encountered in add
    self.W += -learning_rate*grad

lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.371694 val accuracy: 0.388000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.362755 val accuracy: 0.368000
lr 5.000000e-05 reg 2.500000e+04 train accuracy: 0.112980 val accuracy: 0.120000
lr 5.000000e-05 reg 5.000000e+04 train accuracy: 0.100265 val accuracy: 0.087000
best validation accuracy achieved during cross-validation: 0.388000

```

```

[ ]: # Visualize the cross-validation results
import math
import pdb


# pdb.set_trace()

x_scatter = [math.log10(x[0]) for x in results]
y_scatter = [math.log10(x[1]) for x in results]

# plot training accuracy
marker_size = 100
colors = [results[x][0] for x in results]
plt.subplot(2, 1, 1)
plt.tight_layout(pad=3)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap=plt.cm.coolwarm)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 training accuracy')

# plot validation accuracy
colors = [results[x][1] for x in results] # default size of markers is 20
plt.subplot(2, 1, 2)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap=plt.cm.coolwarm)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 validation accuracy')
plt.show()

```



```
[ ]: # Evaluate the best svm on test set
y_test_pred = best_svm.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)
```

linear SVM on raw pixels final test set accuracy: 0.103000

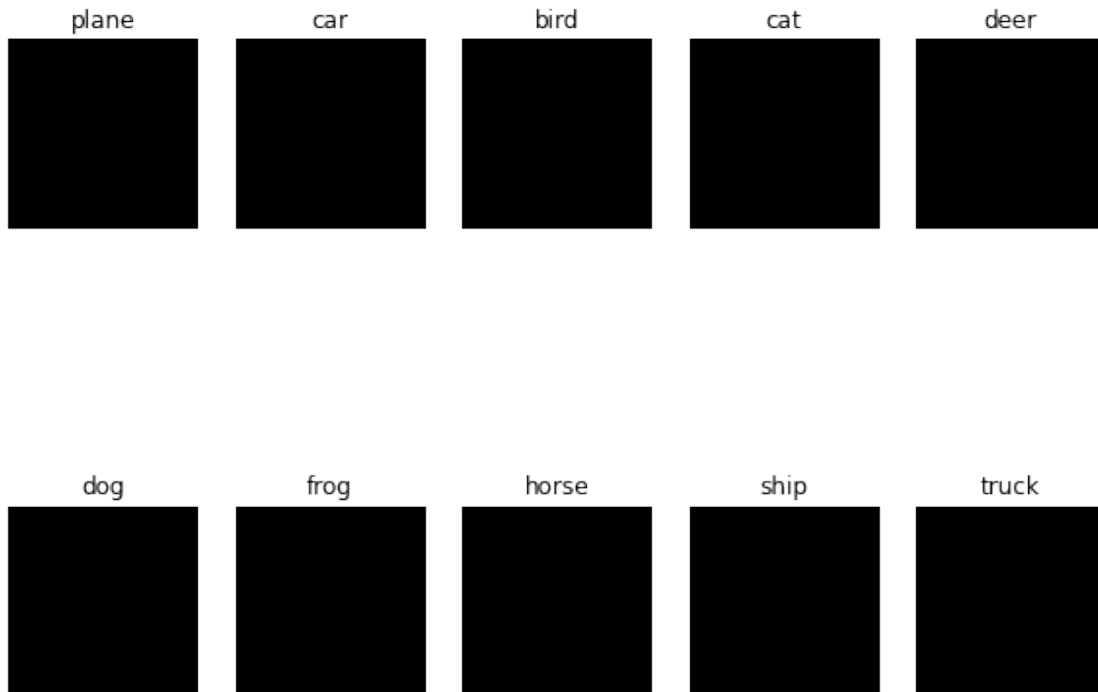
```
[ ]: # Visualize the learned weights for each class.
# Depending on your choice of learning rate and regularization strength, these
    ↪ may
# or may not be nice to look at.
w = best_svm.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)
w_min, w_max = np.min(w), np.max(w)
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
    ↪ 'ship', 'truck']
```

```

for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])

```



### Inline question 2

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way that they do.

*Your Answer :* I seem to have made an error in the code somewhere, my classes images aren't showing up.

*In general, images serve as templates for each class. And the most occurring pixel properties (features) in each image has a big role in deciding the weights for the whole class. Since most images of a same class are more or less similar, these most occurring pixels decide the weights for the class.*

# softmax

October 9, 2022

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'ENPM809K Assignments/assignment1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
↳ -----
↳
↳ MessageError                                Traceback (most recent call↳
↳ last)
↳
↳ <ipython-input-2-4cbfb96d0dd9> in <module>
↳     1 # This mounts your Google Drive to the Colab VM.
↳     2 from google.colab import drive
↳ ----> 3 drive.mount('/content/drive', force_remount=True)
↳     4
↳     5 # Enter the foldername in your Drive where you have saved the↳
↳ unzipped
```

```

/usr/local/lib/python3.7/dist-packages/google/colab/drive.py in
↳ mount(mountpoint, force_remount, timeout_ms)
    103         force_remount=force_remount,
    104         timeout_ms=timeout_ms,
--> 105         ephemeral=True)
    106
    107

/usr/local/lib/python3.7/dist-packages/google/colab/drive.py in
↳ _mount(mountpoint, force_remount, timeout_ms, ephemeral)
    118     if ephemeral:
    119         _message.blocking_request(
--> 120             'request_auth', request={'authType': 'dfs_ephemeral'},
↳ timeout_sec=None)
    121
    122     mountpoint = _os.path.expanduser(mountpoint)

/usr/local/lib/python3.7/dist-packages/google/colab/_message.py in
↳ blocking_request(request_type, request, timeout_sec, parent)
    169     request_id = send_request(
    170         request_type, request, parent=parent, expect_reply=True)
--> 171     return read_reply_from_input(request_id, timeout_sec)

/usr/local/lib/python3.7/dist-packages/google/colab/_message.py in
↳ read_reply_from_input(message_id, timeout_sec)
    100         reply.get('colab_msg_id') == message_id):
    101         if 'error' in reply:
--> 102             raise MessageError(reply['error'])
    103         return reply.get('data', None)
    104

```

MessageError: Error: credential propagation was unsuccessful

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

```

↳ -----

MessageError                                Traceback (most recent call↳
↳ last)

```



```

<ipython-input-3-d5df0069828e> in <module>
      1 from google.colab import drive
----> 2 drive.mount('/content/drive')

/usr/local/lib/python3.7/dist-packages/google/colab/drive.py in
↳ mount(mountpoint, force_remount, timeout_ms)
      103         force_remount=force_remount,
      104         timeout_ms=timeout_ms,
--> 105         ephemeral=True)
      106
      107

/usr/local/lib/python3.7/dist-packages/google/colab/drive.py in
↳ _mount(mountpoint, force_remount, timeout_ms, ephemeral)
      118     if ephemeral:
      119         _message.blocking_request(
--> 120             'request_auth', request={'authType': 'dfs_ephemeral'},
↳ timeout_sec=None)
      121
      122     mountpoint = _os.path.expanduser(mountpoint)

/usr/local/lib/python3.7/dist-packages/google/colab/_message.py in
↳ blocking_request(request_type, request, timeout_sec, parent)
      169     request_id = send_request(
      170         request_type, request, parent=parent, expect_reply=True)
--> 171     return read_reply_from_input(request_id, timeout_sec)

/usr/local/lib/python3.7/dist-packages/google/colab/_message.py in
↳ read_reply_from_input(message_id, timeout_sec)
      100         reply.get('colab_msg_id') == message_id):
      101         if 'error' in reply:
--> 102             raise MessageError(reply['error'])
      103         return reply.get('data', None)
      104

```

MessageError: Error: credential propagation was unsuccessful

# 1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization strength**
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[ ]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

[ ]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000,
    → num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    # Cleaning up variables to prevent loading data multiple times (which may
    → cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
```

```

pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# subsample the data
mask = list(range(num_training, num_training + num_validation))
X_val = X_train[mask]
y_val = y_train[mask]
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = □
    ↪ get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)

```

```

print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)

```

```

Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)

```

## 1.1 Softmax Classifier

Your code for this section will all be written inside `cs231n/classifiers/softmax.py`.

```

[ ]: # First implement the naive softmax loss function with nested loops.
     # Open the file cs231n/classifiers/softmax.py and implement the
     # softmax_loss_naive function.

from cs231n.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))

```

```

loss: 2.402285
sanity check: 2.302585

```

### Inline Question 1

Why do we expect our loss to be close to  $-\log(0.1)$ ? Explain briefly.\*\*

*Your Answer :* First of all, if you mathematically calculate the -ve natural log of 0.1 i.e. you get 2.302585093 which is very close to our loss estimated above.

Theoretically speaking, since we have ten classes and initially each class is assigned the same probability, let's say we get exp of correct class =  $x$  and so the sum of all exponents of all classes will be  $10x$ . If you take

$-1.\log(x/10x)$

You get  $-\log(1/10) \sim -\log(0.1)$

Hence our class is closer to that value,  $-\log(0.1)$

```
[ ]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: -1.076924 analytic: -1.076924, relative error: 1.507795e-08
numerical: 2.852019 analytic: 2.852019, relative error: 1.973547e-08
numerical: 0.670762 analytic: 0.670762, relative error: 8.812854e-08
numerical: 0.341616 analytic: 0.341616, relative error: 1.181130e-07
numerical: -2.419663 analytic: -2.419663, relative error: 1.586656e-08
numerical: -1.527610 analytic: -1.527610, relative error: 9.787021e-09
numerical: -0.873218 analytic: -0.873218, relative error: 2.137063e-08
numerical: -1.243857 analytic: -1.243857, relative error: 2.581577e-08
numerical: 1.561009 analytic: 1.561008, relative error: 6.467397e-08
numerical: 3.639652 analytic: 3.639651, relative error: 2.157088e-08
numerical: -0.189145 analytic: -0.189145, relative error: 1.077258e-07
numerical: -1.661756 analytic: -1.661756, relative error: 4.182324e-09
numerical: -0.009731 analytic: -0.009731, relative error: 2.541629e-06
numerical: 0.636686 analytic: 0.636686, relative error: 4.939228e-08
numerical: 1.944569 analytic: 1.944569, relative error: 1.902823e-08
numerical: 1.247422 analytic: 1.247422, relative error: 2.564956e-08
numerical: -2.653095 analytic: -2.653095, relative error: 4.969549e-09
numerical: 0.269826 analytic: 0.269826, relative error: 1.624245e-07
numerical: 1.654755 analytic: 1.654755, relative error: 4.962747e-08
numerical: 4.325124 analytic: 4.325124, relative error: 2.312031e-08
```

```
[ ]: # Now that we have a naive implementation of the softmax loss function and its
    ↪ gradient,
# implement a vectorized version in softmax_loss_vectorized.
# The two versions should compute the same results, but the vectorized version
    ↪ should be
# much faster.
tic = time.time()
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))
```

```

from cs231n.classifiers.softmax import softmax_loss_vectorized
tic = time.time()
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.
    ↳0.000005)
toc = time.time()
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# As we did for the SVM, we use the Frobenius norm to compare the two versions
# of the gradient.
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
print('Gradient difference: %f' % grad_difference)

```

```

naive loss: 2.402285e+00 computed in 0.115508s
vectorized loss: 2.402285e+00 computed in 0.011926s
Loss difference: 0.000000
Gradient difference: 0.000000

```

```

[ ]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.

from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained softmax classifier in best_softmax. #
#####

# Provided as a reference. You may or may not want to change these
    ↳hyperparameters
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e4, 5e4]

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

soft_max = Softmax() #Creating a Softmax instance to access the classifier

#First, we iterate through each learning rate
#Second, for each learning rate, we iterate through each regularization to find
    ↳the fit

```

```

for learn_rate in learning_rates:
    for reg_strength in regularization_strengths:

        #Third, we train our classifier
        loss_function = soft_max.train(X_train, y_train, learning_rate =
        ↪learn_rate, reg = reg_strength, num_iters=1000)

        #Fourth, we predict the classes on training data and estimate the accuracy
        y_pred_train = soft_max.predict(X_train)
        training_accuracy = np.mean(y_pred_train == y_train)

        #Fifth, we predict the classes on the validation set and estimate its
        ↪accuracy
        y_pred_validation = soft_max.predict(X_val)
        validation_accuracy = np.mean(y_pred_validation == y_val)

        #Next, store these accuracies as dictionaries in the result dict.
        results[(learn_rate, reg_strength)] = (training_accuracy,
        ↪validation_accuracy)

        #Lastly, we check if the validation accuracy is better than the best
        ↪validation accuracy and then update it accordingly
        if validation_accuracy > best_val:
            best_val = validation_accuracy
            best_softmax = soft_max #The current softmax will be the best softmax for
            ↪this iteration

"""Note here, that by increasing/decreasing the number of iteration in our
    ↪training of the SVM
we can have an direct effect on our validation accuracy"""

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
    ↪best_val)

```

```

lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.333061 val accuracy: 0.348000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.312898 val accuracy: 0.324000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.332551 val accuracy: 0.339000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.307816 val accuracy: 0.320000
best validation accuracy achieved during cross-validation: 0.348000

```

```
[ ]: # evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))
```

softmax on raw pixels final test set accuracy: 0.326000

### Inline Question 2 - True or False

Suppose the overall training loss is defined as the sum of the per-datapoint loss over all training examples. It is possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

*Your Answer : True*

*Your Explanation : In an SVM classifier, the loss function only serves to check if the loss values are within a certain margin. Once the margins are reached, the SVM classifier is satisfied with its loss value and doesn't delve any deeper. Any new datapoint addition will have no impact on the loss if the margin is satisfied.*

*But with a Softmax classifier, it tries to keep attaining higher probabilities for the correct class and lower probabilities for the wrong class and thus there is always room to lower the costs even lower. Thus any addition of a new data point will have an effect on the Softmax loss.*

```
[ ]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
↳ 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```





[ ]:

## two\_layer\_net

October 9, 2022

```
[2]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'ENPM809K Assignments/assignment1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Mounted at /content/drive

/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/datasets

/content/drive/My Drive/ENPM809K Assignments/assignment1

## 1 Fully-Connected Neural Nets

In this exercise we will implement fully-connected networks using a modular approach. For each layer we will implement a **forward** and a **backward** function. The **forward** function will receive inputs, weights, and other parameters and will return both an output and a **cache** object storing data needed for the backward pass, like this:

```
def layer_forward(x, w):
    """ Receive inputs x and weights w """
    # Do some computations ...
    z = # ... some intermediate value
    # Do some more computations ...
```

```

out = # the output

cache = (x, w, z, out) # Values we need to compute gradients

return out, cache

```

The backward pass will receive upstream derivatives and the `cache` object, and will return gradients with respect to the inputs and weights, like this:

```

def layer_backward(dout, cache):
    """
    Receive dout (derivative of loss with respect to outputs) and cache,
    and compute derivative with respect to inputs.
    """
    # Unpack cache values
    x, w, z, out = cache

    # Use values in cache to compute derivatives
    dx = # Derivative of loss with respect to x
    dw = # Derivative of loss with respect to w

    return dx, dw

```

After implementing a bunch of layers this way, we will be able to easily combine them to build classifiers with different architectures.

```

[40]: # As usual, a bit of setup
from __future__ import print_function
import time
import numpy as np
import matplotlib.pyplot as plt
from cs231n.classifiers.fc_net import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, \
    eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):

```

```

""" returns relative error """
return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

[4]: *# Load the (preprocessed) CIFAR10 data.*

```

data = get_CIFAR10_data()
for k, v in list(data.items()):
    print('%s: ' % k, v.shape)

('X_train: ', (49000, 3, 32, 32))
('y_train: ', (49000,))
('X_val: ', (1000, 3, 32, 32))
('y_val: ', (1000,))
('X_test: ', (1000, 3, 32, 32))
('y_test: ', (1000,))

```

## 2 Affine layer: forward

Open the file `cs231n/layers.py` and implement the `affine_forward` function.

Once you are done you can test your implementation by running the following:

[5]: *# Test the affine\_forward function*

```

num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape),
→output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,  3.77273342]])

# Compare your output with ours. The error should be around e-9 or less.
print('Testing affine_forward function:')
print('difference: ', rel_error(out, correct_out))

```

Testing affine\_forward function:  
difference: 9.769849468192957e-10

### 3 Affine layer: backward

Now implement the `affine_backward` function and test your implementation using numeric gradient checking.

```
[6]: # Test the affine_backward function
np.random.seed(231)
x = np.random.randn(10, 2, 3)
w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x,
    ↪dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w,
    ↪dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b,
    ↪dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be around e-10 or less
print('Testing affine_backward function:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing affine\_backward function:  
dx error: 5.399100368651805e-11  
dw error: 9.904211865398145e-11  
db error: 2.4122867568119087e-11

### 4 ReLU activation: forward

Implement the forward pass for the ReLU activation function in the `relu_forward` function and test your implementation using the following:

```
[7]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)
```

```

out, _ = relu_forward(x)
correct_out = np.array([[ 0.,          0.,          0.,          0.,          ],
                        [ 0.,          0.,          0.04545455, 0.13636364,],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5,          ]])

# Compare your output with ours. The error should be on the order of e-8
print('Testing relu_forward function:')
print('difference: ', rel_error(out, correct_out))

```

Testing relu\_forward function:  
difference: 4.999999798022158e-08

## 5 ReLU activation: backward

Now implement the backward pass for the ReLU activation function in the `relu_backward` function and test your implementation using numeric gradient checking:

```

[8]: np.random.seed(231)
x = np.random.randn(10, 10)
dout = np.random.randn(*x.shape)

dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

_, cache = relu_forward(x)
dx = relu_backward(dout, cache)

# The error should be on the order of e-12
print('Testing relu_backward function:')
print('dx error: ', rel_error(dx_num, dx))

```

Testing relu\_backward function:  
dx error: 3.2756349136310288e-12

### 5.1 Inline Question 1:

We've only asked you to implement ReLU, but there are a number of different activation functions that one could use in neural networks, each with its pros and cons. In particular, an issue commonly seen with activation functions is getting zero (or close to zero) gradient flow during backpropagation. Which of the following activation functions have this problem? If you consider these functions in the one dimensional case, what types of input would lead to this behaviour? 1. Sigmoid 2. ReLU 3. Leaky ReLU

## 5.2 Answer: 1 and 2

1. Sigmoids saturate and kill gradients. The sigmoid activation pushes the gradient to zero when the input values for it are closer 0 and on the other end of the spectrum, closer to 1
2. ReLU units can be fragile during training and can “die”. During backprop, if large gradients are flowing through, it can cause the weights to be updated such that the unit doesn’t update for newer datapoints from that instance onwards. And that causes your gradient to be set to zero.
3. Leaky ReLUs are designed so that they can handle the “dying” of the ReLU units. Hence, they do not have this problem.

## 6 “Sandwich” layers

There are some common patterns of layers that are frequently used in neural nets. For example, affine layers are frequently followed by a ReLU nonlinearity. To make these common patterns easy, we define several convenience layers in the file `cs231n/layer_utils.py`.

For now take a look at the `affine_relu_forward` and `affine_relu_backward` functions, and run the following to numerically gradient check the backward pass:

```
[9]: from cs231n.layer_utils import affine_relu_forward, affine_relu_backward
np.random.seed(231)
x = np.random.randn(2, 3, 4)
w = np.random.randn(12, 10)
b = np.random.randn(10)
dout = np.random.randn(2, 10)

out, cache = affine_relu_forward(x, w, b)
dx, dw, db = affine_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w, b)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w, b)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w, b)[0], b, dout)

# Relative error should be around e-10 or less
print('Testing affine_relu_forward and affine_relu_backward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing affine\_relu\_forward and affine\_relu\_backward:

```
dx error:  2.299579177309368e-11
dw error:  8.162011105764925e-11
db error:  7.826724021458994e-12
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## 7 Loss layers: Softmax and SVM

Now implement the loss and gradient for softmax and SVM in the `softmax_loss` and `svm_loss` function in `cs231n/layers.py`. These should be similar to what you implemented in `cs231n/classifiers/softmax.py` and `cs231n/classifiers/linear_svm.py`.

You can make sure that the implementations are correct by running the following:

```
[15]: np.random.seed(231)
num_classes, num_inputs = 10, 50
x = 0.001 * np.random.randn(num_inputs, num_classes)
y = np.random.randint(num_classes, size=num_inputs)

dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
loss, dx = svm_loss(x, y)

# Test svm_loss function. Loss should be around 9 and dx error should be around
# the order of e-9
print('Testing svm_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x,
# verbose=False)
loss, dx = softmax_loss(x, y)

# Test softmax_loss function. Loss should be close to 2.3 and dx error should
# be around e-8
print('\nTesting softmax_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))
```

```
Testing svm_loss:
loss: 8.999602749096233
dx error: 1.4021566006651672e-09
```

```
Testing softmax_loss:
loss: 2.3025458445007376
dx error: 8.234144091578429e-09
```



## 8 Two-layer network

Open the file `cs231n/classifiers/fc_net.py` and complete the implementation of the `TwoLayerNet` class. Read through it to make sure you understand the API. You can run the cell below to test your implementation.

```
[20]: np.random.seed(231)
N, D, H, C = 3, 5, 50, 7
X = np.random.randn(N, D)
y = np.random.randint(C, size=N)

std = 1e-3
model = TwoLayerNet(input_dim=D, hidden_dim=H, num_classes=C, weight_scale=std)

print('Testing initialization ... ')
W1_std = abs(model.params['W1'].std() - std)
b1 = model.params['b1']
W2_std = abs(model.params['W2'].std() - std)
b2 = model.params['b2']
assert W1_std < std / 10, 'First layer weights do not seem right'
assert np.all(b1 == 0), 'First layer biases do not seem right'
assert W2_std < std / 10, 'Second layer weights do not seem right'
assert np.all(b2 == 0), 'Second layer biases do not seem right'

print('Testing test-time forward pass ... ')
model.params['W1'] = np.linspace(-0.7, 0.3, num=D*H).reshape(D, H)
model.params['b1'] = np.linspace(-0.1, 0.9, num=H)
model.params['W2'] = np.linspace(-0.3, 0.4, num=H*C).reshape(H, C)
model.params['b2'] = np.linspace(-0.9, 0.1, num=C)
X = np.linspace(-5.5, 4.5, num=N*D).reshape(D, N).T
scores = model.loss(X)
correct_scores = np.asarray(
    [[11.53165108, 12.2917344, 13.05181771, 13.81190102, 14.57198434, 15.
↪ 33206765, 16.09215096],
    [12.05769098, 12.74614105, 13.43459113, 14.1230412, 14.81149128, 15.
↪ 49994135, 16.18839143],
    [12.58373087, 13.20054771, 13.81736455, 14.43418138, 15.05099822, 15.
↪ 66781506, 16.2846319 ]])
scores_diff = np.abs(scores - correct_scores).sum()
assert scores_diff < 1e-6, 'Problem with test-time forward pass'

print('Testing training loss (no regularization)')
y = np.asarray([0, 5, 1])
loss, grads = model.loss(X, y)
correct_loss = 3.4702243556
assert abs(loss - correct_loss) < 1e-10, 'Problem with training-time loss'
```

```

model.reg = 1.0
loss, grads = model.loss(X, y)
correct_loss = 26.5948426952
assert abs(loss - correct_loss) < 1e-10, 'Problem with regularization loss'

# Errors should be around e-7 or less
for reg in [0.0, 0.7]:
    print('Running numeric gradient check with reg = ', reg)
    model.reg = reg
    loss, grads = model.loss(X, y)

    for name in sorted(grads):
        f = lambda _: model.loss(X, y)[0]
        grad_num = eval_numerical_gradient(f, model.params[name], verbose=False)
        print('%s relative error: %.2e' % (name, rel_error(grad_num, grads[name])))

```

```

Testing initialization ...
Testing test-time forward pass ...
Testing training loss (no regularization)
Running numeric gradient check with reg = 0.0
W1 relative error: 1.83e-08
W2 relative error: 3.31e-10
b1 relative error: 9.83e-09
b2 relative error: 4.33e-10
Running numeric gradient check with reg = 0.7
W1 relative error: 2.53e-07
W2 relative error: 2.85e-08
b1 relative error: 1.56e-08
b2 relative error: 7.76e-10

```

## 9 Solver

Open the file `cs231n/solver.py` and read through it to familiarize yourself with the API. You also need to implement the `sgd` function in `cs231n/optim.py`. After doing so, use a `Solver` instance to train a `TwoLayerNet` that achieves about 36% accuracy on the validation set.

```

[41]: input_size = 32 * 32 * 3
      hidden_size = 50
      num_classes = 10
      model = TwoLayerNet(input_size, hidden_size, num_classes)
      solver = None

#####
# TODO: Use a Solver instance to train a TwoLayerNet that achieves about 36% #
# accuracy on the validation set.                                           #
#####

```

```

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

#Here we use the same syntax to create a solver instance on the model
solver = Solver(model, data,
                update_rule='sgd',
                optim_config={
                    'learning_rate': 1e-3,
                },
                lr_decay=0.95,
                num_epochs=10, batch_size=100,
                print_every=100)

solver.train()

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####

```

```

(Iteration 1 / 4900) loss: 2.303803
(Epoch 0 / 10) train acc: 0.160000; val_acc: 0.176000
(Iteration 101 / 4900) loss: 1.909482
(Iteration 201 / 4900) loss: 1.717844
(Iteration 301 / 4900) loss: 1.661954
(Iteration 401 / 4900) loss: 1.798607
(Epoch 1 / 10) train acc: 0.423000; val_acc: 0.422000
(Iteration 501 / 4900) loss: 1.641987
(Iteration 601 / 4900) loss: 1.554380
(Iteration 701 / 4900) loss: 1.443837
(Iteration 801 / 4900) loss: 1.536855
(Iteration 901 / 4900) loss: 1.360803
(Epoch 2 / 10) train acc: 0.483000; val_acc: 0.457000
(Iteration 1001 / 4900) loss: 1.563823
(Iteration 1101 / 4900) loss: 1.551614
(Iteration 1201 / 4900) loss: 1.390415
(Iteration 1301 / 4900) loss: 1.624959
(Iteration 1401 / 4900) loss: 1.399313
(Epoch 3 / 10) train acc: 0.468000; val_acc: 0.454000
(Iteration 1501 / 4900) loss: 1.491097
(Iteration 1601 / 4900) loss: 1.515698
(Iteration 1701 / 4900) loss: 1.571268
(Iteration 1801 / 4900) loss: 1.402826
(Iteration 1901 / 4900) loss: 1.337489
(Epoch 4 / 10) train acc: 0.511000; val_acc: 0.477000
(Iteration 2001 / 4900) loss: 1.474201
(Iteration 2101 / 4900) loss: 1.175097
(Iteration 2201 / 4900) loss: 1.280630
(Iteration 2301 / 4900) loss: 1.389493

```

```
(Iteration 2401 / 4900) loss: 1.250720
(Epoch 5 / 10) train acc: 0.524000; val_acc: 0.492000
(Iteration 2501 / 4900) loss: 1.366097
(Iteration 2601 / 4900) loss: 1.412278
(Iteration 2701 / 4900) loss: 1.445154
(Iteration 2801 / 4900) loss: 1.339657
(Iteration 2901 / 4900) loss: 1.465240
(Epoch 6 / 10) train acc: 0.524000; val_acc: 0.501000
(Iteration 3001 / 4900) loss: 1.071470
(Iteration 3101 / 4900) loss: 1.199761
(Iteration 3201 / 4900) loss: 1.375785
(Iteration 3301 / 4900) loss: 1.179885
(Iteration 3401 / 4900) loss: 1.479837
(Epoch 7 / 10) train acc: 0.552000; val_acc: 0.493000
(Iteration 3501 / 4900) loss: 1.474006
(Iteration 3601 / 4900) loss: 1.148525
(Iteration 3701 / 4900) loss: 1.233577
(Iteration 3801 / 4900) loss: 1.194387
(Iteration 3901 / 4900) loss: 1.423228
(Epoch 8 / 10) train acc: 0.544000; val_acc: 0.480000
(Iteration 4001 / 4900) loss: 1.371165
(Iteration 4101 / 4900) loss: 1.205742
(Iteration 4201 / 4900) loss: 1.133436
(Iteration 4301 / 4900) loss: 1.260861
(Iteration 4401 / 4900) loss: 1.069663
(Epoch 9 / 10) train acc: 0.561000; val_acc: 0.501000
(Iteration 4501 / 4900) loss: 1.240798
(Iteration 4601 / 4900) loss: 1.211501
(Iteration 4701 / 4900) loss: 1.329369
(Iteration 4801 / 4900) loss: 1.259118
(Epoch 10 / 10) train acc: 0.580000; val_acc: 0.497000
```

## 10 Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.36 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
[42]: # Run this cell to visualize training loss and train / val accuracy
```


```
plt.subplot(2, 1, 1)
```

```

plt.title('Training loss')
plt.plot(solver.loss_history, 'o')
plt.xlabel('Iteration')

plt.subplot(2, 1, 2)
plt.title('Accuracy')
plt.plot(solver.train_acc_history, '-o', label='train')
plt.plot(solver.val_acc_history, '-o', label='val')
plt.plot([0.5] * len(solver.val_acc_history), 'k--')
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.gcf().set_size_inches(15, 12)
plt.show()

```



```


[43]: from cs231n.vis_utils import visualize_grid

      # Visualize the weights of the network

```

```
def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.reshape(32, 32, 3, -1).transpose(3, 0, 1, 2)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```



## 11 Tune your hyperparameters

**What's wrong?.** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the learning rate decay, but you should be able to get good performance using the default value.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 48% on the validation set. Our best network gets over 52% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can (52% could serve as a reference), with a fully-connected Neural Network. Feel free implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

```
[44]: best_model = None

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
  ↪ #
# model in best_model.
  ↪ #
#
  ↪ #
# To help debug your network, it may help to use visualizations similar to the
  ↪ #
# ones we used above; these visualizations will have significant qualitative
  ↪ #
# differences from the ones we saw above for the poorly tuned network.
  ↪ #
#
  ↪ #
# Tweaking hyperparameters by hand can be fun, but you might find it useful to
  ↪ #
# write code to sweep through possible combinations of hyperparameters
  ↪ #
# automatically like we did on thes previous exercises.
  ↪ #
```

```
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

results = {}
best_val = -1

#sources for how to arrive at a safe range of values for learning rate and
↳regularization strengths
#https://machinelearningmastery.com/
↳learning-rate-for-deep-learning-neural-networks/
#https://cs231n.github.io/neural-networks-3/#hyper
learning_rates = [1e-7, 1e-5, 1e-3, 1e-1]
regularization_strengths = [1e-7, 1e-5, 1e-3, 1e-1]

count_learn = 0 #just to checkout iterations

#First, we iterate through each learning rate
#Second, for each learning rate, we iterate through each regularization to find
↳the fit
for learn_rate in learning_rates:
    count_reg = 0
    for reg_strength in regularization_strengths:
        print("Iteration: lr_iteration = %d reg_iteration = %d" %(count_learn,
↳count_reg))
        #Third, we create our model
        model = TwoLayerNet(reg = reg_strength)
        #Fourth, we make a solver instance for the model
        solver = Solver(model, data,
                        update_rule='sgd',
                        optim_config={
                            'learning_rate': learn_rate,
                        },
                        lr_decay=0.95,
                        num_epochs=4, batch_size=100,
                        print_every=100)
        #Fifth, we train the model
        solver.train()
        #Sixth, store the best validation accuracy as dictionaries in the result
↳dict for the current values of lr and reg.
        results[(learn_rate, reg_strength)] = solver.best_val_acc

        #Lastly, we check if the validation accuracy is better than the best
↳validation accuracy and then update it accordingly
        if results[(learn_rate, reg_strength)] > best_val:
            best_val = results[(learn_rate, reg_strength)]
```



```

        best_model = model #The current model will be the best model for this
        ↪iteration
        count_reg += 1 #increasing out counter just understand the iterations
        count_learn += 1

# Print out results.
for lr, reg in sorted(results):
    val_accuracy = results[(lr, reg)]
    print('lr %e reg %e val accuracy: %f' % (
        lr, reg, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
    ↪best_val)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####

```

```

Iteration: lr_iteration = 0 reg_iteration = 0
(Iteration 1 / 1960) loss: 2.301694
(Epoch 0 / 4) train acc: 0.091000; val_acc: 0.087000
(Iteration 101 / 1960) loss: 2.301338
(Iteration 201 / 1960) loss: 2.303732
(Iteration 301 / 1960) loss: 2.300110
(Iteration 401 / 1960) loss: 2.304532
(Epoch 1 / 4) train acc: 0.095000; val_acc: 0.090000
(Iteration 501 / 1960) loss: 2.302802
(Iteration 601 / 1960) loss: 2.303712
(Iteration 701 / 1960) loss: 2.303146
(Iteration 801 / 1960) loss: 2.299986
(Iteration 901 / 1960) loss: 2.304796
(Epoch 2 / 4) train acc: 0.091000; val_acc: 0.094000
(Iteration 1001 / 1960) loss: 2.304845
(Iteration 1101 / 1960) loss: 2.303961
(Iteration 1201 / 1960) loss: 2.306147
(Iteration 1301 / 1960) loss: 2.307568
(Iteration 1401 / 1960) loss: 2.296721
(Epoch 3 / 4) train acc: 0.078000; val_acc: 0.095000
(Iteration 1501 / 1960) loss: 2.302441
(Iteration 1601 / 1960) loss: 2.300383
(Iteration 1701 / 1960) loss: 2.305257
(Iteration 1801 / 1960) loss: 2.299982
(Iteration 1901 / 1960) loss: 2.301426
(Epoch 4 / 4) train acc: 0.073000; val_acc: 0.099000
Iteration: lr_iteration = 0 reg_iteration = 1

```

```

(Iteration 1 / 1960) loss: 2.301907
(Epoch 0 / 4) train acc: 0.143000; val_acc: 0.114000
(Iteration 101 / 1960) loss: 2.302682
(Iteration 201 / 1960) loss: 2.301838
(Iteration 301 / 1960) loss: 2.303224
(Iteration 401 / 1960) loss: 2.300711
(Epoch 1 / 4) train acc: 0.123000; val_acc: 0.115000
(Iteration 501 / 1960) loss: 2.298307
(Iteration 601 / 1960) loss: 2.299892
(Iteration 701 / 1960) loss: 2.302852
(Iteration 801 / 1960) loss: 2.303961
(Iteration 901 / 1960) loss: 2.301457
(Epoch 2 / 4) train acc: 0.117000; val_acc: 0.116000
(Iteration 1001 / 1960) loss: 2.301302
(Iteration 1101 / 1960) loss: 2.303048
(Iteration 1201 / 1960) loss: 2.301878
(Iteration 1301 / 1960) loss: 2.298768
(Iteration 1401 / 1960) loss: 2.298694
(Epoch 3 / 4) train acc: 0.140000; val_acc: 0.121000
(Iteration 1501 / 1960) loss: 2.303858
(Iteration 1601 / 1960) loss: 2.296444
(Iteration 1701 / 1960) loss: 2.303409
(Iteration 1801 / 1960) loss: 2.304978
(Iteration 1901 / 1960) loss: 2.296995
(Epoch 4 / 4) train acc: 0.111000; val_acc: 0.123000
Iteration: lr_iteration = 0 reg_iteration = 2
(Iteration 1 / 1960) loss: 2.303242
(Epoch 0 / 4) train acc: 0.089000; val_acc: 0.094000
(Iteration 101 / 1960) loss: 2.305584
(Iteration 201 / 1960) loss: 2.304883
(Iteration 301 / 1960) loss: 2.307286
(Iteration 401 / 1960) loss: 2.303852
(Epoch 1 / 4) train acc: 0.092000; val_acc: 0.094000
(Iteration 501 / 1960) loss: 2.301540
(Iteration 601 / 1960) loss: 2.304709
(Iteration 701 / 1960) loss: 2.305134
(Iteration 801 / 1960) loss: 2.306481
(Iteration 901 / 1960) loss: 2.302322
(Epoch 2 / 4) train acc: 0.081000; val_acc: 0.096000
(Iteration 1001 / 1960) loss: 2.301762
(Iteration 1101 / 1960) loss: 2.308629
(Iteration 1201 / 1960) loss: 2.301431
(Iteration 1301 / 1960) loss: 2.306605
(Iteration 1401 / 1960) loss: 2.303644
(Epoch 3 / 4) train acc: 0.080000; val_acc: 0.095000
(Iteration 1501 / 1960) loss: 2.302910
(Iteration 1601 / 1960) loss: 2.303939
(Iteration 1701 / 1960) loss: 2.305173

```

(Iteration 1801 / 1960) loss: 2.302236  
(Iteration 1901 / 1960) loss: 2.306249  
(Epoch 4 / 4) train acc: 0.115000; val\_acc: 0.095000  
Iteration: lr\_iteration = 0 reg\_iteration = 3  
(Iteration 1 / 1960) loss: 2.320262  
(Epoch 0 / 4) train acc: 0.072000; val\_acc: 0.083000  
(Iteration 101 / 1960) loss: 2.318271  
(Iteration 201 / 1960) loss: 2.325128  
(Iteration 301 / 1960) loss: 2.319290  
(Iteration 401 / 1960) loss: 2.317636  
(Epoch 1 / 4) train acc: 0.093000; val\_acc: 0.085000  
(Iteration 501 / 1960) loss: 2.319644  
(Iteration 601 / 1960) loss: 2.320207  
(Iteration 701 / 1960) loss: 2.315540  
(Iteration 801 / 1960) loss: 2.318976  
(Iteration 901 / 1960) loss: 2.317351  
(Epoch 2 / 4) train acc: 0.088000; val\_acc: 0.087000  
(Iteration 1001 / 1960) loss: 2.316805  
(Iteration 1101 / 1960) loss: 2.317758  
(Iteration 1201 / 1960) loss: 2.317393  
(Iteration 1301 / 1960) loss: 2.317627  
(Iteration 1401 / 1960) loss: 2.319093  
(Epoch 3 / 4) train acc: 0.076000; val\_acc: 0.089000  
(Iteration 1501 / 1960) loss: 2.315982  
(Iteration 1601 / 1960) loss: 2.319944  
(Iteration 1701 / 1960) loss: 2.319648  
(Iteration 1801 / 1960) loss: 2.316362  
(Iteration 1901 / 1960) loss: 2.314719  
(Epoch 4 / 4) train acc: 0.106000; val\_acc: 0.091000  
Iteration: lr\_iteration = 1 reg\_iteration = 0  
(Iteration 1 / 1960) loss: 2.301595  
(Epoch 0 / 4) train acc: 0.112000; val\_acc: 0.137000  
(Iteration 101 / 1960) loss: 2.291074  
(Iteration 201 / 1960) loss: 2.288734  
(Iteration 301 / 1960) loss: 2.275866  
(Iteration 401 / 1960) loss: 2.278741  
(Epoch 1 / 4) train acc: 0.200000; val\_acc: 0.228000  
(Iteration 501 / 1960) loss: 2.257521  
(Iteration 601 / 1960) loss: 2.262762  
(Iteration 701 / 1960) loss: 2.249515  
(Iteration 801 / 1960) loss: 2.248811  
(Iteration 901 / 1960) loss: 2.241182  
(Epoch 2 / 4) train acc: 0.199000; val\_acc: 0.239000  
(Iteration 1001 / 1960) loss: 2.223598  
(Iteration 1101 / 1960) loss: 2.224300  
(Iteration 1201 / 1960) loss: 2.197657  
(Iteration 1301 / 1960) loss: 2.208012  
(Iteration 1401 / 1960) loss: 2.197264

(Epoch 3 / 4) train acc: 0.251000; val\_acc: 0.240000  
(Iteration 1501 / 1960) loss: 2.190154  
(Iteration 1601 / 1960) loss: 2.158686  
(Iteration 1701 / 1960) loss: 2.152834  
(Iteration 1801 / 1960) loss: 2.198179  
(Iteration 1901 / 1960) loss: 2.133406  
(Epoch 4 / 4) train acc: 0.244000; val\_acc: 0.264000  
Iteration: lr\_iteration = 1 reg\_iteration = 1  
(Iteration 1 / 1960) loss: 2.300059  
(Epoch 0 / 4) train acc: 0.098000; val\_acc: 0.100000  
(Iteration 101 / 1960) loss: 2.298942  
(Iteration 201 / 1960) loss: 2.290234  
(Iteration 301 / 1960) loss: 2.285474  
(Iteration 401 / 1960) loss: 2.284032  
(Epoch 1 / 4) train acc: 0.205000; val\_acc: 0.191000  
(Iteration 501 / 1960) loss: 2.268767  
(Iteration 601 / 1960) loss: 2.267911  
(Iteration 701 / 1960) loss: 2.258763  
(Iteration 801 / 1960) loss: 2.250399  
(Iteration 901 / 1960) loss: 2.230070  
(Epoch 2 / 4) train acc: 0.218000; val\_acc: 0.225000  
(Iteration 1001 / 1960) loss: 2.242796  
(Iteration 1101 / 1960) loss: 2.207299  
(Iteration 1201 / 1960) loss: 2.241002  
(Iteration 1301 / 1960) loss: 2.209093  
(Iteration 1401 / 1960) loss: 2.212782  
(Epoch 3 / 4) train acc: 0.227000; val\_acc: 0.243000  
(Iteration 1501 / 1960) loss: 2.185756  
(Iteration 1601 / 1960) loss: 2.190009  
(Iteration 1701 / 1960) loss: 2.162732  
(Iteration 1801 / 1960) loss: 2.126370  
(Iteration 1901 / 1960) loss: 2.099603  
(Epoch 4 / 4) train acc: 0.239000; val\_acc: 0.256000  
Iteration: lr\_iteration = 1 reg\_iteration = 2  
(Iteration 1 / 1960) loss: 2.304034  
(Epoch 0 / 4) train acc: 0.087000; val\_acc: 0.087000  
(Iteration 101 / 1960) loss: 2.302844  
(Iteration 201 / 1960) loss: 2.293354  
(Iteration 301 / 1960) loss: 2.288782  
(Iteration 401 / 1960) loss: 2.289444  
(Epoch 1 / 4) train acc: 0.221000; val\_acc: 0.234000  
(Iteration 501 / 1960) loss: 2.273328  
(Iteration 601 / 1960) loss: 2.268822  
(Iteration 701 / 1960) loss: 2.264653  
(Iteration 801 / 1960) loss: 2.252770  
(Iteration 901 / 1960) loss: 2.255248  
(Epoch 2 / 4) train acc: 0.235000; val\_acc: 0.237000  
(Iteration 1001 / 1960) loss: 2.245632

```

(Iteration 1101 / 1960) loss: 2.203940
(Iteration 1201 / 1960) loss: 2.218811
(Iteration 1301 / 1960) loss: 2.223688
(Iteration 1401 / 1960) loss: 2.204227
(Epoch 3 / 4) train acc: 0.217000; val_acc: 0.232000
(Iteration 1501 / 1960) loss: 2.159685
(Iteration 1601 / 1960) loss: 2.183794
(Iteration 1701 / 1960) loss: 2.195502
(Iteration 1801 / 1960) loss: 2.187506
(Iteration 1901 / 1960) loss: 2.135558
(Epoch 4 / 4) train acc: 0.226000; val_acc: 0.240000
Iteration: lr_iteration = 1 reg_iteration = 3
(Iteration 1 / 1960) loss: 2.316607
(Epoch 0 / 4) train acc: 0.096000; val_acc: 0.087000
(Iteration 101 / 1960) loss: 2.311923
(Iteration 201 / 1960) loss: 2.309658
(Iteration 301 / 1960) loss: 2.302567
(Iteration 401 / 1960) loss: 2.295370
(Epoch 1 / 4) train acc: 0.217000; val_acc: 0.216000
(Iteration 501 / 1960) loss: 2.288181
(Iteration 601 / 1960) loss: 2.291940
(Iteration 701 / 1960) loss: 2.271879
(Iteration 801 / 1960) loss: 2.266277
(Iteration 901 / 1960) loss: 2.280460
(Epoch 2 / 4) train acc: 0.218000; val_acc: 0.232000
(Iteration 1001 / 1960) loss: 2.243532
(Iteration 1101 / 1960) loss: 2.219654
(Iteration 1201 / 1960) loss: 2.253645
(Iteration 1301 / 1960) loss: 2.249395
(Iteration 1401 / 1960) loss: 2.229151
(Epoch 3 / 4) train acc: 0.217000; val_acc: 0.248000
(Iteration 1501 / 1960) loss: 2.158345
(Iteration 1601 / 1960) loss: 2.226519
(Iteration 1701 / 1960) loss: 2.199641
(Iteration 1801 / 1960) loss: 2.171442
(Iteration 1901 / 1960) loss: 2.076453
(Epoch 4 / 4) train acc: 0.218000; val_acc: 0.256000
Iteration: lr_iteration = 2 reg_iteration = 0
(Iteration 1 / 1960) loss: 2.300819
(Epoch 0 / 4) train acc: 0.139000; val_acc: 0.133000
(Iteration 101 / 1960) loss: 1.832876
(Iteration 201 / 1960) loss: 1.820312
(Iteration 301 / 1960) loss: 1.887772
(Iteration 401 / 1960) loss: 1.568764
(Epoch 1 / 4) train acc: 0.446000; val_acc: 0.447000
(Iteration 501 / 1960) loss: 1.626059
(Iteration 601 / 1960) loss: 1.477519
(Iteration 701 / 1960) loss: 1.480743

```

(Iteration 801 / 1960) loss: 1.538570  
(Iteration 901 / 1960) loss: 1.642470  
(Epoch 2 / 4) train acc: 0.526000; val\_acc: 0.482000  
(Iteration 1001 / 1960) loss: 1.496164  
(Iteration 1101 / 1960) loss: 1.343786  
(Iteration 1201 / 1960) loss: 1.212589  
(Iteration 1301 / 1960) loss: 1.236717  
(Iteration 1401 / 1960) loss: 1.364237  
(Epoch 3 / 4) train acc: 0.515000; val\_acc: 0.484000  
(Iteration 1501 / 1960) loss: 1.230296  
(Iteration 1601 / 1960) loss: 1.349812  
(Iteration 1701 / 1960) loss: 1.359621  
(Iteration 1801 / 1960) loss: 1.366731  
(Iteration 1901 / 1960) loss: 1.300014  
(Epoch 4 / 4) train acc: 0.508000; val\_acc: 0.485000  
Iteration: lr\_iteration = 2 reg\_iteration = 1  
(Iteration 1 / 1960) loss: 2.304827  
(Epoch 0 / 4) train acc: 0.143000; val\_acc: 0.159000  
(Iteration 101 / 1960) loss: 1.828599  
(Iteration 201 / 1960) loss: 1.854298  
(Iteration 301 / 1960) loss: 1.642512  
(Iteration 401 / 1960) loss: 1.713940  
(Epoch 1 / 4) train acc: 0.434000; val\_acc: 0.434000  
(Iteration 501 / 1960) loss: 1.564445  
(Iteration 601 / 1960) loss: 1.374773  
(Iteration 701 / 1960) loss: 1.586927  
(Iteration 801 / 1960) loss: 1.468098  
(Iteration 901 / 1960) loss: 1.526729  
(Epoch 2 / 4) train acc: 0.481000; val\_acc: 0.466000  
(Iteration 1001 / 1960) loss: 1.316097  
(Iteration 1101 / 1960) loss: 1.461357  
(Iteration 1201 / 1960) loss: 1.626224  
(Iteration 1301 / 1960) loss: 1.795160  
(Iteration 1401 / 1960) loss: 1.295571  
(Epoch 3 / 4) train acc: 0.503000; val\_acc: 0.486000  
(Iteration 1501 / 1960) loss: 1.386841  
(Iteration 1601 / 1960) loss: 1.386159  
(Iteration 1701 / 1960) loss: 1.290159  
(Iteration 1801 / 1960) loss: 1.514143  
(Iteration 1901 / 1960) loss: 1.074882  
(Epoch 4 / 4) train acc: 0.491000; val\_acc: 0.508000  
Iteration: lr\_iteration = 2 reg\_iteration = 2  
(Iteration 1 / 1960) loss: 2.301756  
(Epoch 0 / 4) train acc: 0.135000; val\_acc: 0.125000  
(Iteration 101 / 1960) loss: 1.725061  
(Iteration 201 / 1960) loss: 1.801686  
(Iteration 301 / 1960) loss: 1.642149  
(Iteration 401 / 1960) loss: 1.640869

```

(Epoch 1 / 4) train acc: 0.460000; val_acc: 0.435000
(Iteration 501 / 1960) loss: 1.373048
(Iteration 601 / 1960) loss: 1.345676
(Iteration 701 / 1960) loss: 1.381678
(Iteration 801 / 1960) loss: 1.312563
(Iteration 901 / 1960) loss: 1.624729
(Epoch 2 / 4) train acc: 0.492000; val_acc: 0.461000
(Iteration 1001 / 1960) loss: 1.489512
(Iteration 1101 / 1960) loss: 1.631681
(Iteration 1201 / 1960) loss: 1.671602
(Iteration 1301 / 1960) loss: 1.302067
(Iteration 1401 / 1960) loss: 1.384110
(Epoch 3 / 4) train acc: 0.517000; val_acc: 0.477000
(Iteration 1501 / 1960) loss: 1.434011
(Iteration 1601 / 1960) loss: 1.369640
(Iteration 1701 / 1960) loss: 1.319401
(Iteration 1801 / 1960) loss: 1.596330
(Iteration 1901 / 1960) loss: 1.504938
(Epoch 4 / 4) train acc: 0.533000; val_acc: 0.487000
Iteration: lr_iteration = 2 reg_iteration = 3
(Iteration 1 / 1960) loss: 2.316900
(Epoch 0 / 4) train acc: 0.156000; val_acc: 0.166000
(Iteration 101 / 1960) loss: 1.710005
(Iteration 201 / 1960) loss: 1.563495
(Iteration 301 / 1960) loss: 1.598511
(Iteration 401 / 1960) loss: 1.257848
(Epoch 1 / 4) train acc: 0.441000; val_acc: 0.454000
(Iteration 501 / 1960) loss: 1.594175
(Iteration 601 / 1960) loss: 1.498441
(Iteration 701 / 1960) loss: 1.350564
(Iteration 801 / 1960) loss: 1.580846
(Iteration 901 / 1960) loss: 1.496661
(Epoch 2 / 4) train acc: 0.461000; val_acc: 0.474000
(Iteration 1001 / 1960) loss: 1.850866
(Iteration 1101 / 1960) loss: 1.493007
(Iteration 1201 / 1960) loss: 1.421010
(Iteration 1301 / 1960) loss: 1.406938
(Iteration 1401 / 1960) loss: 1.603426
(Epoch 3 / 4) train acc: 0.475000; val_acc: 0.463000
(Iteration 1501 / 1960) loss: 1.407619
(Iteration 1601 / 1960) loss: 1.386524
(Iteration 1701 / 1960) loss: 1.407189
(Iteration 1801 / 1960) loss: 1.362849
(Iteration 1901 / 1960) loss: 1.406931
(Epoch 4 / 4) train acc: 0.541000; val_acc: 0.510000
Iteration: lr_iteration = 3 reg_iteration = 0
(Iteration 1 / 1960) loss: 2.304093
(Epoch 0 / 4) train acc: 0.132000; val_acc: 0.125000

```

```

/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/layers.py:247:
RuntimeWarning: invalid value encountered in true_divide
/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/layers.py:249:
RuntimeWarning: divide by zero encountered in log
/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/layers.py:255:
RuntimeWarning: invalid value encountered in true_divide
/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/layers.py:257:
RuntimeWarning: invalid value encountered in true_divide

(Iteration 101 / 1960) loss: nan
(Iteration 201 / 1960) loss: nan
(Iteration 301 / 1960) loss: nan
(Iteration 401 / 1960) loss: nan
(Epoch 1 / 4) train acc: 0.105000; val_acc: 0.087000
(Iteration 501 / 1960) loss: nan
(Iteration 601 / 1960) loss: nan
(Iteration 701 / 1960) loss: nan
(Iteration 801 / 1960) loss: nan
(Iteration 901 / 1960) loss: nan
(Epoch 2 / 4) train acc: 0.105000; val_acc: 0.087000
(Iteration 1001 / 1960) loss: nan
(Iteration 1101 / 1960) loss: nan
(Iteration 1201 / 1960) loss: nan
(Iteration 1301 / 1960) loss: nan
(Iteration 1401 / 1960) loss: nan
(Epoch 3 / 4) train acc: 0.096000; val_acc: 0.087000
(Iteration 1501 / 1960) loss: nan
(Iteration 1601 / 1960) loss: nan
(Iteration 1701 / 1960) loss: nan
(Iteration 1801 / 1960) loss: nan
(Iteration 1901 / 1960) loss: nan
(Epoch 4 / 4) train acc: 0.096000; val_acc: 0.087000
Iteration: lr_iteration = 3 reg_iteration = 1
(Iteration 1 / 1960) loss: 2.300467
(Epoch 0 / 4) train acc: 0.095000; val_acc: 0.081000
(Iteration 101 / 1960) loss: nan
(Iteration 201 / 1960) loss: nan
(Iteration 301 / 1960) loss: nan
(Iteration 401 / 1960) loss: nan
(Epoch 1 / 4) train acc: 0.109000; val_acc: 0.087000
(Iteration 501 / 1960) loss: nan
(Iteration 601 / 1960) loss: nan
(Iteration 701 / 1960) loss: nan
(Iteration 801 / 1960) loss: nan
(Iteration 901 / 1960) loss: nan
(Epoch 2 / 4) train acc: 0.087000; val_acc: 0.087000
(Iteration 1001 / 1960) loss: nan
(Iteration 1101 / 1960) loss: nan

```



```

(Iteration 1201 / 1960) loss: nan
(Iteration 1301 / 1960) loss: nan
(Iteration 1401 / 1960) loss: nan
(Epoch 3 / 4) train acc: 0.102000; val_acc: 0.087000
(Iteration 1501 / 1960) loss: nan
(Iteration 1601 / 1960) loss: nan
(Iteration 1701 / 1960) loss: nan
(Iteration 1801 / 1960) loss: nan
(Iteration 1901 / 1960) loss: nan
(Epoch 4 / 4) train acc: 0.121000; val_acc: 0.087000
Iteration: lr_iteration = 3 reg_iteration = 2
(Iteration 1 / 1960) loss: 2.305783
(Epoch 0 / 4) train acc: 0.098000; val_acc: 0.090000
(Iteration 101 / 1960) loss: nan
(Iteration 201 / 1960) loss: nan
(Iteration 301 / 1960) loss: nan
(Iteration 401 / 1960) loss: nan
(Epoch 1 / 4) train acc: 0.107000; val_acc: 0.087000
(Iteration 501 / 1960) loss: nan
(Iteration 601 / 1960) loss: nan
(Iteration 701 / 1960) loss: nan
(Iteration 801 / 1960) loss: nan
(Iteration 901 / 1960) loss: nan
(Epoch 2 / 4) train acc: 0.097000; val_acc: 0.087000
(Iteration 1001 / 1960) loss: nan
(Iteration 1101 / 1960) loss: nan
(Iteration 1201 / 1960) loss: nan
(Iteration 1301 / 1960) loss: nan
(Iteration 1401 / 1960) loss: nan
(Epoch 3 / 4) train acc: 0.092000; val_acc: 0.087000
(Iteration 1501 / 1960) loss: nan
(Iteration 1601 / 1960) loss: nan
(Iteration 1701 / 1960) loss: nan
(Iteration 1801 / 1960) loss: nan
(Iteration 1901 / 1960) loss: nan
(Epoch 4 / 4) train acc: 0.116000; val_acc: 0.087000
Iteration: lr_iteration = 3 reg_iteration = 3
(Iteration 1 / 1960) loss: 2.318338
(Epoch 0 / 4) train acc: 0.125000; val_acc: 0.137000
(Iteration 101 / 1960) loss: nan
(Iteration 201 / 1960) loss: nan
(Iteration 301 / 1960) loss: nan
(Iteration 401 / 1960) loss: nan
(Epoch 1 / 4) train acc: 0.102000; val_acc: 0.087000
(Iteration 501 / 1960) loss: nan
(Iteration 601 / 1960) loss: nan
(Iteration 701 / 1960) loss: nan
(Iteration 801 / 1960) loss: nan

```

```

(Iteration 901 / 1960) loss: nan
(Epoch 2 / 4) train acc: 0.095000; val_acc: 0.087000
(Iteration 1001 / 1960) loss: nan
(Iteration 1101 / 1960) loss: nan
(Iteration 1201 / 1960) loss: nan
(Iteration 1301 / 1960) loss: nan
(Iteration 1401 / 1960) loss: nan
(Epoch 3 / 4) train acc: 0.107000; val_acc: 0.087000
(Iteration 1501 / 1960) loss: nan
(Iteration 1601 / 1960) loss: nan
(Iteration 1701 / 1960) loss: nan
(Iteration 1801 / 1960) loss: nan
(Iteration 1901 / 1960) loss: nan
(Epoch 4 / 4) train acc: 0.088000; val_acc: 0.087000
lr 1.000000e-07 reg 1.000000e-07 val accuracy: 0.099000
lr 1.000000e-07 reg 1.000000e-05 val accuracy: 0.123000
lr 1.000000e-07 reg 1.000000e-03 val accuracy: 0.096000
lr 1.000000e-07 reg 1.000000e-01 val accuracy: 0.091000
lr 1.000000e-05 reg 1.000000e-07 val accuracy: 0.264000
lr 1.000000e-05 reg 1.000000e-05 val accuracy: 0.256000
lr 1.000000e-05 reg 1.000000e-03 val accuracy: 0.240000
lr 1.000000e-05 reg 1.000000e-01 val accuracy: 0.256000
lr 1.000000e-03 reg 1.000000e-07 val accuracy: 0.485000
lr 1.000000e-03 reg 1.000000e-05 val accuracy: 0.508000
lr 1.000000e-03 reg 1.000000e-03 val accuracy: 0.487000
lr 1.000000e-03 reg 1.000000e-01 val accuracy: 0.510000
lr 1.000000e-01 reg 1.000000e-07 val accuracy: 0.125000
lr 1.000000e-01 reg 1.000000e-05 val accuracy: 0.087000
lr 1.000000e-01 reg 1.000000e-03 val accuracy: 0.090000
lr 1.000000e-01 reg 1.000000e-01 val accuracy: 0.137000
best validation accuracy achieved during cross-validation: 0.510000

```

## 12 Test your model!

Run your best model on the validation and test sets. You should achieve above 48% accuracy on the validation set and the test set.

```

[45]: y_val_pred = np.argmax(best_model.loss(data['X_val']), axis=1)
      print('Validation set accuracy: ', (y_val_pred == data['y_val']).mean())

```

Validation set accuracy: 0.51

```

[46]: y_test_pred = np.argmax(best_model.loss(data['X_test']), axis=1)
      print('Test set accuracy: ', (y_test_pred == data['y_test']).mean())

```

Test set accuracy: 0.505

## 12.1 Inline Question 2:

Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

\$Your Answer:\$1 and 3

*Your Explanation :* 1. By increasing the size of the dataset you are providing more varied dataset to your model and if its is split proportionally between your train and test, you will improve the test accuracy as it has more data refine its parameters such as weights and biases etc. 2. Adding more hidden units will not really affect the magnitude of the accuracy because both the test and train data will pass through the same ones. 3. Regularization strength pushes the model to prevent overfitting and arriving at the best model possible which mking sure to not let the model get complex. Increasing this will only push the model to increase its efficiency, thus increasing the efficiency of the testing set.

[ ]:

# features

October 9, 2022

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'ENPM809K Assignments/assignment1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My Drive/ENPM809K Assignments/assignment1/cs231n/datasets
/content/drive/My Drive/ENPM809K Assignments/assignment1
```

## 1 Image features exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

We have seen that we can achieve reasonable performance on an image classification task by training a linear classifier on the pixels of the input image. In this exercise we will show that we can improve our classification performance by training linear classifiers not on raw pixels but on features that are computed from the raw pixels.

All of your work for this exercise will be done in this notebook.

```
[2]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
→ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

## 1.1 Load data

Similar to previous exercises, we will load CIFAR-10 data from disk.

```
[3]: from cs231n.features import color_histogram_hsv, hog_feature

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    # Cleaning up variables to prevent loading data multiple times (which may
    → cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
        pass

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
```

```

X_test = X_test[mask]
y_test = y_test[mask]

return X_train, y_train, X_val, y_val, X_test, y_test

X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()

```

## 1.2 Extract Features

For each image we will compute a Histogram of Oriented Gradients (HOG) as well as a color histogram using the hue channel in HSV color space. We form our final feature vector for each image by concatenating the HOG and color histogram feature vectors.

Roughly speaking, HOG should capture the texture of the image while ignoring color information, and the color histogram represents the color of the input image while ignoring texture. As a result, we expect that using both together ought to work better than using either alone. Verifying this assumption would be a good thing to try for your own interest.

The `hog_feature` and `color_histogram_hsv` functions both operate on a single image and return a feature vector for that image. The `extract_features` function takes a set of images and a list of feature functions and evaluates each feature function on each image, storing the results in a matrix where each column is the concatenation of all feature vectors for a single image.

```

[4]: from cs231n.features import *

num_color_bins = 10 # Number of bins in the color histogram
feature_fns = [hog_feature, lambda img: color_histogram_hsv(img,
    ↪nbin=num_color_bins)]
X_train_feats = extract_features(X_train, feature_fns, verbose=True)
X_val_feats = extract_features(X_val, feature_fns)
X_test_feats = extract_features(X_test, feature_fns)

# Preprocessing: Subtract the mean feature
mean_feat = np.mean(X_train_feats, axis=0, keepdims=True)
X_train_feats -= mean_feat
X_val_feats -= mean_feat
X_test_feats -= mean_feat

# Preprocessing: Divide by standard deviation. This ensures that each feature
# has roughly the same scale.
std_feat = np.std(X_train_feats, axis=0, keepdims=True)
X_train_feats /= std_feat
X_val_feats /= std_feat
X_test_feats /= std_feat

# Preprocessing: Add a bias dimension
X_train_feats = np.hstack([X_train_feats, np.ones((X_train_feats.shape[0], 1))])

```

```
X_val_feats = np.hstack([X_val_feats, np.ones((X_val_feats.shape[0], 1))])
X_test_feats = np.hstack([X_test_feats, np.ones((X_test_feats.shape[0], 1))])
```

```
Done extracting features for 1000 / 49000 images
Done extracting features for 2000 / 49000 images
Done extracting features for 3000 / 49000 images
Done extracting features for 4000 / 49000 images
Done extracting features for 5000 / 49000 images
Done extracting features for 6000 / 49000 images
Done extracting features for 7000 / 49000 images
Done extracting features for 8000 / 49000 images
Done extracting features for 9000 / 49000 images
Done extracting features for 10000 / 49000 images
Done extracting features for 11000 / 49000 images
Done extracting features for 12000 / 49000 images
Done extracting features for 13000 / 49000 images
Done extracting features for 14000 / 49000 images
Done extracting features for 15000 / 49000 images
Done extracting features for 16000 / 49000 images
Done extracting features for 17000 / 49000 images
Done extracting features for 18000 / 49000 images
Done extracting features for 19000 / 49000 images
Done extracting features for 20000 / 49000 images
Done extracting features for 21000 / 49000 images
Done extracting features for 22000 / 49000 images
Done extracting features for 23000 / 49000 images
Done extracting features for 24000 / 49000 images
Done extracting features for 25000 / 49000 images
Done extracting features for 26000 / 49000 images
Done extracting features for 27000 / 49000 images
Done extracting features for 28000 / 49000 images
Done extracting features for 29000 / 49000 images
Done extracting features for 30000 / 49000 images
Done extracting features for 31000 / 49000 images
Done extracting features for 32000 / 49000 images
Done extracting features for 33000 / 49000 images
Done extracting features for 34000 / 49000 images
Done extracting features for 35000 / 49000 images
Done extracting features for 36000 / 49000 images
Done extracting features for 37000 / 49000 images
Done extracting features for 38000 / 49000 images
Done extracting features for 39000 / 49000 images
Done extracting features for 40000 / 49000 images
Done extracting features for 41000 / 49000 images
Done extracting features for 42000 / 49000 images
Done extracting features for 43000 / 49000 images
Done extracting features for 44000 / 49000 images
```

```

Done extracting features for 45000 / 49000 images
Done extracting features for 46000 / 49000 images
Done extracting features for 47000 / 49000 images
Done extracting features for 48000 / 49000 images
Done extracting features for 49000 / 49000 images

```

### 1.3 Train SVM on features

Using the multiclass SVM code developed earlier in the assignment, train SVMs on top of the features extracted above; this should achieve better results than training SVMs directly on top of raw pixels.

```

[5]: # Use the validation set to tune the learning rate and regularization strength

from cs231n.classifiers.linear_classifier import LinearSVM

#using the same learning rates and regularization strengths as last time
learning_rates = [1e-7, 1e-5, 1e-3, 1e-1]
regularization_strengths = [1e-7, 1e-5, 1e-3, 1e-1]

results = {}
best_val = -1
best_svm = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained classifier in best_svm. You might also want to play #
# with different numbers of bins in the color histogram. If you are careful #
# you should be able to get accuracy of near 0.44 on the validation set. #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

svm = LinearSVM() #Creating an svm instance to access the classifier

#First, we iterate through each learning rate
#Second, for each learning rate, we iterate through each regularization to find
    ↳ the fit
for learn_rate in learning_rates:
    for reg_strength in regularization_strengths:

        #Third, we train our classifier, but instead of training it on all image
        ↳ pixels, we'll do it on image features
        loss_function = svm.train(X_train_feats, y_train, learning_rate =
        ↳ learn_rate, reg = reg_strength, num_iters=1000)

```



```

    #Fourth, we predict the classes on training data's features and estimate
    ↳ the accuracy
    y_pred_train = svm.predict(X_train_feats)
    training_accuracy = np.mean(y_pred_train == y_train)

    #Fifth, we predict the classes on the validation set's features and
    ↳ estimate its accuracy
    y_pred_validation = svm.predict(X_val_feats)
    validation_accuracy = np.mean(y_pred_validation == y_val)

    #Next, store these accuracies as dictionaries in the result dict.
    results[(learn_rate, reg_strength)] = (training_accuracy,
    ↳ validation_accuracy)

    #Lastly, we check if the validation accuracy is better than the best
    ↳ validation accuracy and then update it accordingly
    if validation_accuracy > best_val:
        best_val = validation_accuracy
        best_svm = svm #The current SVM will be the best SVM for this iteration

    """Note here, that by increasing/decreasing the number of iteration in our
    ↳ training of the SVM
    we can have an direct effect on our validation accuracy"""

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved: %f' % best_val)

```

```

lr 1.000000e-07 reg 1.000000e-07 train accuracy: 0.130755 val accuracy: 0.149000
lr 1.000000e-07 reg 1.000000e-05 train accuracy: 0.143286 val accuracy: 0.152000
lr 1.000000e-07 reg 1.000000e-03 train accuracy: 0.155082 val accuracy: 0.161000
lr 1.000000e-07 reg 1.000000e-01 train accuracy: 0.166673 val accuracy: 0.167000
lr 1.000000e-05 reg 1.000000e-07 train accuracy: 0.409959 val accuracy: 0.413000
lr 1.000000e-05 reg 1.000000e-05 train accuracy: 0.415347 val accuracy: 0.420000
lr 1.000000e-05 reg 1.000000e-03 train accuracy: 0.416531 val accuracy: 0.424000
lr 1.000000e-05 reg 1.000000e-01 train accuracy: 0.417878 val accuracy: 0.429000
lr 1.000000e-03 reg 1.000000e-07 train accuracy: 0.496612 val accuracy: 0.484000
lr 1.000000e-03 reg 1.000000e-05 train accuracy: 0.505143 val accuracy: 0.494000
lr 1.000000e-03 reg 1.000000e-03 train accuracy: 0.508816 val accuracy: 0.489000
lr 1.000000e-03 reg 1.000000e-01 train accuracy: 0.510245 val accuracy: 0.489000
lr 1.000000e-01 reg 1.000000e-07 train accuracy: 0.484776 val accuracy: 0.469000

```

```

lr 1.000000e-01 reg 1.000000e-05 train accuracy: 0.481673 val accuracy: 0.489000
lr 1.000000e-01 reg 1.000000e-03 train accuracy: 0.485673 val accuracy: 0.479000
lr 1.000000e-01 reg 1.000000e-01 train accuracy: 0.467367 val accuracy: 0.438000
best validation accuracy achieved: 0.494000

```

```

[6]: # Evaluate your trained SVM on the test set: you should be able to get at least
      ↪0.40
      y_test_pred = best_svm.predict(X_test_feats)
      test_accuracy = np.mean(y_test == y_test_pred)
      print(test_accuracy)

```


0.468

```

[7]: # An important way to gain intuition about how an algorithm works is to
      # visualize the mistakes that it makes. In this visualization, we show examples
      # of images that are misclassified by our current system. The first column
      # shows images that our system labeled as "plane" but whose true label is
      # something other than "plane".

      examples_per_class = 8
      classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
                  ↪'ship', 'truck']
      for cls, cls_name in enumerate(classes):
          idxs = np.where((y_test != cls) & (y_test_pred == cls))[0]
          idxs = np.random.choice(idxs, examples_per_class, replace=False)
          for i, idx in enumerate(idxs):
              plt.subplot(examples_per_class, len(classes), i * len(classes) + cls +
                  ↪1)
              plt.imshow(X_test[idx].astype('uint8'))
              plt.axis('off')
              if i == 0:
                  plt.title(cls_name)
      plt.show()

```



### 1.3.1 Inline question 1:

Describe the misclassification results that you see. Do they make sense?

*Your Answer :* There is definitely a scope for misclassification. Coming to describing them, if you see them on a higher level, it seems like it doesn't make any sense. But upon a closer look, since we are concatenating two different kinds of features to arrive at a feature vector, namely the Histogram of Gradients (HoG) and Colour Histogram. With HoG, it classifies images with the similar features and their orientations. But since we are not actually measuring pixel to pixel, it doesn't really take into account where that oriented gradient is, it just sees that two images have an equal distribution of the oriented gradients in its histogram, so the classifier assumes that both images must be from the same class. Similarly, with the histogram of colours, if between two images the colour distribution histogram is more or less the same, the images are classified as a same class.

In my classifier, it makes sense at some places and not much sense with others. It has classified a lot of trucks as cars. Since both of them, more or less, would have similar looking features (body, wheels, tyres etc.), their HoG might have been such that the trucks would've been classified as cars. Even in the case of the planes class, it is simply classing wrong images which might pass

of as a plane (i.e too much blue in the background) as the plane class (maybe the result of the features produced from the colour Histogram).

Now in some cases, it doesn't make sense as to why a certain image was classified as such, for eg. a cat was classified as a plane. Upon studying the image, one can't really point as to which features of the cat must've contributed to the model for it to be classified as a plane. In such cases, the prediction is just bad and one can't take away a lot from it.

Note: The example above answer may not exactly match the image result produced, because with each compilation, a different set of images are produced. I saw this change after having written the original answer.

## 1.4 Neural Network on image features

Earlier in this assignment we saw that training a two-layer neural network on raw pixels achieved better classification performance than linear classifiers on raw pixels. In this notebook we have seen that linear classifiers on image features outperform linear classifiers on raw pixels.

For completeness, we should also try training a neural network on image features. This approach should outperform all previous approaches: you should easily be able to achieve over 55% classification accuracy on the test set; our best model achieves about 60% classification accuracy.

```
[8]: # Preprocessing: Remove the bias dimension
# Make sure to run this cell only ONCE
print(X_train_feats.shape)
X_train_feats = X_train_feats[:, :-1]
X_val_feats = X_val_feats[:, :-1]
X_test_feats = X_test_feats[:, :-1]

print(X_train_feats.shape)
```

```
(49000, 155)
```

```
(49000, 154)
```

```
[9]: from cs231n.classifiers.fc_net import TwoLayerNet
from cs231n.solver import Solver

input_dim = X_train_feats.shape[1]
hidden_dim = 600 #tweaked this to get required results
num_classes = 10

best_net = None

#####
# TODO: Train a two-layer neural network on image features. You may want to #
# cross-validate various parameters as in previous sections. Store your best #
# model in the best_net variable. #
```

```
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

results = {}
best_val = -1

#same sets for learning rates and regularization strengths as last time
learning_rates = [1e-7, 1e-5, 1e-3, 1e-1]
regularization_strengths = [1e-7, 1e-5, 1e-3, 1e-1]

#saving the input features as train, validation, and test
#Saving their respective labels in the y variables
data = {
    'X_train': X_train_feats, # training data
    'y_train': y_train, # training labels
    'X_val': X_val_feats, # validation data
    'y_val': y_val, # validation labels
    'X_test': X_test_feats, #test data
    'y_test': y_test, #test labels
}

count_learn = 0 #just used to understand the iterations in the output
#First, we iterate through each learning rate
#Second, for each learning rate, we iterate through each regularization to find
    ↳the fit
for learn_rate in learning_rates:
    count_reg = 0
    for reg_strength in regularization_strengths:
        #Third, we create our model
        print("Iteration: lr_iteration = %d reg_iteration = %d" %(count_learn,
    ↳count_reg))
        net = TwoLayerNet(input_dim, hidden_dim, num_classes, reg= reg_strength)
        #Fourth, we make a solver instance for the model
        solver = Solver(net, data,
                        update_rule='sgd',
                        optim_config={
                            'learning_rate': learn_rate,
                        },
                        lr_decay=0.95,
                        num_epochs=4, batch_size=100,
                        print_every=100)
        #Fifth, we train the model
        solver.train()
        #Sixth, store the best validation accuracy as dictionaries in the result
    ↳dict for the current values of lr and reg.
        results[(learn_rate, reg_strength)] = solver.best_val_acc
```

```

    #Lastly, we check if the validation accuracy is better than the best
    ↪validation accuracy and then update it accordingly
    if results[(learn_rate, reg_strength)] > best_val:
        best_val = results[(learn_rate, reg_strength)]
        best_net = net #The current model will be the best model for this
    ↪iteration
        count_reg += 1 #increasing out counter just understand the iterations
    count_learn += 1

# Print out results.
for lr, reg in sorted(results):
    val_accuracy = results[(lr, reg)]
    print('lr %e reg %e val accuracy: %f' % (
        lr, reg, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
    ↪best_val)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

```

```

Iteration: lr_iteration = 0 reg_iteration = 0
(Iteration 1 / 1960) loss: 2.302644
(Epoch 0 / 4) train acc: 0.101000; val_acc: 0.096000
(Iteration 101 / 1960) loss: 2.302611
(Iteration 201 / 1960) loss: 2.302603
(Iteration 301 / 1960) loss: 2.302601
(Iteration 401 / 1960) loss: 2.302618
(Epoch 1 / 4) train acc: 0.098000; val_acc: 0.096000
(Iteration 501 / 1960) loss: 2.302597
(Iteration 601 / 1960) loss: 2.302569
(Iteration 701 / 1960) loss: 2.302556
(Iteration 801 / 1960) loss: 2.302566
(Iteration 901 / 1960) loss: 2.302607
(Epoch 2 / 4) train acc: 0.105000; val_acc: 0.096000
(Iteration 1001 / 1960) loss: 2.302594
(Iteration 1101 / 1960) loss: 2.302599
(Iteration 1201 / 1960) loss: 2.302587
(Iteration 1301 / 1960) loss: 2.302594
(Iteration 1401 / 1960) loss: 2.302589
(Epoch 3 / 4) train acc: 0.099000; val_acc: 0.096000
(Iteration 1501 / 1960) loss: 2.302579
(Iteration 1601 / 1960) loss: 2.302596
(Iteration 1701 / 1960) loss: 2.302584
(Iteration 1801 / 1960) loss: 2.302646
(Iteration 1901 / 1960) loss: 2.302596

```

```

(Epoch 4 / 4) train acc: 0.095000; val_acc: 0.096000
Iteration: lr_iteration = 0 reg_iteration = 1
(Iteration 1 / 1960) loss: 2.302577
(Epoch 0 / 4) train acc: 0.086000; val_acc: 0.089000
(Iteration 101 / 1960) loss: 2.302571
(Iteration 201 / 1960) loss: 2.302574
(Iteration 301 / 1960) loss: 2.302577
(Iteration 401 / 1960) loss: 2.302599
(Epoch 1 / 4) train acc: 0.104000; val_acc: 0.089000
(Iteration 501 / 1960) loss: 2.302628
(Iteration 601 / 1960) loss: 2.302598
(Iteration 701 / 1960) loss: 2.302560
(Iteration 801 / 1960) loss: 2.302592
(Iteration 901 / 1960) loss: 2.302606
(Epoch 2 / 4) train acc: 0.090000; val_acc: 0.090000
(Iteration 1001 / 1960) loss: 2.302564
(Iteration 1101 / 1960) loss: 2.302572
(Iteration 1201 / 1960) loss: 2.302579
(Iteration 1301 / 1960) loss: 2.302624
(Iteration 1401 / 1960) loss: 2.302603
(Epoch 3 / 4) train acc: 0.095000; val_acc: 0.089000
(Iteration 1501 / 1960) loss: 2.302581
(Iteration 1601 / 1960) loss: 2.302612
(Iteration 1701 / 1960) loss: 2.302582
(Iteration 1801 / 1960) loss: 2.302587
(Iteration 1901 / 1960) loss: 2.302556
(Epoch 4 / 4) train acc: 0.085000; val_acc: 0.089000
Iteration: lr_iteration = 0 reg_iteration = 2
(Iteration 1 / 1960) loss: 2.302668
(Epoch 0 / 4) train acc: 0.086000; val_acc: 0.093000
(Iteration 101 / 1960) loss: 2.302649
(Iteration 201 / 1960) loss: 2.302671
(Iteration 301 / 1960) loss: 2.302628
(Iteration 401 / 1960) loss: 2.302689
(Epoch 1 / 4) train acc: 0.090000; val_acc: 0.093000
(Iteration 501 / 1960) loss: 2.302645
(Iteration 601 / 1960) loss: 2.302643
(Iteration 701 / 1960) loss: 2.302671
(Iteration 801 / 1960) loss: 2.302631
(Iteration 901 / 1960) loss: 2.302650
(Epoch 2 / 4) train acc: 0.078000; val_acc: 0.093000
(Iteration 1001 / 1960) loss: 2.302627
(Iteration 1101 / 1960) loss: 2.302633
(Iteration 1201 / 1960) loss: 2.302659
(Iteration 1301 / 1960) loss: 2.302640
(Iteration 1401 / 1960) loss: 2.302667
(Epoch 3 / 4) train acc: 0.089000; val_acc: 0.093000
(Iteration 1501 / 1960) loss: 2.302650

```

```

(Iteration 1601 / 1960) loss: 2.302678
(Iteration 1701 / 1960) loss: 2.302668
(Iteration 1801 / 1960) loss: 2.302662
(Iteration 1901 / 1960) loss: 2.302647
(Epoch 4 / 4) train acc: 0.088000; val_acc: 0.093000
Iteration: lr_iteration = 0 reg_iteration = 3
(Iteration 1 / 1960) loss: 2.307503
(Epoch 0 / 4) train acc: 0.091000; val_acc: 0.103000
(Iteration 101 / 1960) loss: 2.307503
(Iteration 201 / 1960) loss: 2.307488
(Iteration 301 / 1960) loss: 2.307499
(Iteration 401 / 1960) loss: 2.307474
(Epoch 1 / 4) train acc: 0.098000; val_acc: 0.103000
(Iteration 501 / 1960) loss: 2.307499
(Iteration 601 / 1960) loss: 2.307530
(Iteration 701 / 1960) loss: 2.307515
(Iteration 801 / 1960) loss: 2.307511
(Iteration 901 / 1960) loss: 2.307490
(Epoch 2 / 4) train acc: 0.081000; val_acc: 0.103000
(Iteration 1001 / 1960) loss: 2.307515
(Iteration 1101 / 1960) loss: 2.307489
(Iteration 1201 / 1960) loss: 2.307511
(Iteration 1301 / 1960) loss: 2.307517
(Iteration 1401 / 1960) loss: 2.307530
(Epoch 3 / 4) train acc: 0.099000; val_acc: 0.103000
(Iteration 1501 / 1960) loss: 2.307500
(Iteration 1601 / 1960) loss: 2.307524
(Iteration 1701 / 1960) loss: 2.307464
(Iteration 1801 / 1960) loss: 2.307514
(Iteration 1901 / 1960) loss: 2.307496
(Epoch 4 / 4) train acc: 0.107000; val_acc: 0.103000
Iteration: lr_iteration = 1 reg_iteration = 0
(Iteration 1 / 1960) loss: 2.302612
(Epoch 0 / 4) train acc: 0.101000; val_acc: 0.073000
(Iteration 101 / 1960) loss: 2.302551
(Iteration 201 / 1960) loss: 2.302615
(Iteration 301 / 1960) loss: 2.302589
(Iteration 401 / 1960) loss: 2.302580
(Epoch 1 / 4) train acc: 0.092000; val_acc: 0.074000
(Iteration 501 / 1960) loss: 2.302620
(Iteration 601 / 1960) loss: 2.302593
(Iteration 701 / 1960) loss: 2.302577
(Iteration 801 / 1960) loss: 2.302617
(Iteration 901 / 1960) loss: 2.302592
(Epoch 2 / 4) train acc: 0.071000; val_acc: 0.074000
(Iteration 1001 / 1960) loss: 2.302593
(Iteration 1101 / 1960) loss: 2.302570
(Iteration 1201 / 1960) loss: 2.302607

```



```

(Iteration 1301 / 1960) loss: 2.302610
(Iteration 1401 / 1960) loss: 2.302601
(Epoch 3 / 4) train acc: 0.097000; val_acc: 0.077000
(Iteration 1501 / 1960) loss: 2.302609
(Iteration 1601 / 1960) loss: 2.302636
(Iteration 1701 / 1960) loss: 2.302584
(Iteration 1801 / 1960) loss: 2.302561
(Iteration 1901 / 1960) loss: 2.302594
(Epoch 4 / 4) train acc: 0.098000; val_acc: 0.077000
Iteration: lr_iteration = 1 reg_iteration = 1
(Iteration 1 / 1960) loss: 2.302597
(Epoch 0 / 4) train acc: 0.087000; val_acc: 0.086000
(Iteration 101 / 1960) loss: 2.302579
(Iteration 201 / 1960) loss: 2.302569
(Iteration 301 / 1960) loss: 2.302579
(Iteration 401 / 1960) loss: 2.302600
(Epoch 1 / 4) train acc: 0.097000; val_acc: 0.090000
(Iteration 501 / 1960) loss: 2.302586
(Iteration 601 / 1960) loss: 2.302595
(Iteration 701 / 1960) loss: 2.302596
(Iteration 801 / 1960) loss: 2.302605
(Iteration 901 / 1960) loss: 2.302528
(Epoch 2 / 4) train acc: 0.094000; val_acc: 0.087000
(Iteration 1001 / 1960) loss: 2.302554
(Iteration 1101 / 1960) loss: 2.302560
(Iteration 1201 / 1960) loss: 2.302574
(Iteration 1301 / 1960) loss: 2.302581
(Iteration 1401 / 1960) loss: 2.302573
(Epoch 3 / 4) train acc: 0.089000; val_acc: 0.089000
(Iteration 1501 / 1960) loss: 2.302585
(Iteration 1601 / 1960) loss: 2.302588
(Iteration 1701 / 1960) loss: 2.302577
(Iteration 1801 / 1960) loss: 2.302570
(Iteration 1901 / 1960) loss: 2.302550
(Epoch 4 / 4) train acc: 0.099000; val_acc: 0.089000
Iteration: lr_iteration = 1 reg_iteration = 2
(Iteration 1 / 1960) loss: 2.302668
(Epoch 0 / 4) train acc: 0.102000; val_acc: 0.089000
(Iteration 101 / 1960) loss: 2.302639
(Iteration 201 / 1960) loss: 2.302617
(Iteration 301 / 1960) loss: 2.302629
(Iteration 401 / 1960) loss: 2.302636
(Epoch 1 / 4) train acc: 0.095000; val_acc: 0.094000
(Iteration 501 / 1960) loss: 2.302621
(Iteration 601 / 1960) loss: 2.302650
(Iteration 701 / 1960) loss: 2.302610
(Iteration 801 / 1960) loss: 2.302641
(Iteration 901 / 1960) loss: 2.302616

```

```

(Epoch 2 / 4) train acc: 0.080000; val_acc: 0.094000
(Iteration 1001 / 1960) loss: 2.302658
(Iteration 1101 / 1960) loss: 2.302654
(Iteration 1201 / 1960) loss: 2.302632
(Iteration 1301 / 1960) loss: 2.302654
(Iteration 1401 / 1960) loss: 2.302657
(Epoch 3 / 4) train acc: 0.109000; val_acc: 0.092000
(Iteration 1501 / 1960) loss: 2.302635
(Iteration 1601 / 1960) loss: 2.302641
(Iteration 1701 / 1960) loss: 2.302625
(Iteration 1801 / 1960) loss: 2.302640
(Iteration 1901 / 1960) loss: 2.302627
(Epoch 4 / 4) train acc: 0.092000; val_acc: 0.092000
Iteration: lr_iteration = 1 reg_iteration = 3
(Iteration 1 / 1960) loss: 2.307557
(Epoch 0 / 4) train acc: 0.086000; val_acc: 0.077000
(Iteration 101 / 1960) loss: 2.307487
(Iteration 201 / 1960) loss: 2.307520
(Iteration 301 / 1960) loss: 2.307501
(Iteration 401 / 1960) loss: 2.307491
(Epoch 1 / 4) train acc: 0.088000; val_acc: 0.079000
(Iteration 501 / 1960) loss: 2.307513
(Iteration 601 / 1960) loss: 2.307480
(Iteration 701 / 1960) loss: 2.307490
(Iteration 801 / 1960) loss: 2.307479
(Iteration 901 / 1960) loss: 2.307469
(Epoch 2 / 4) train acc: 0.106000; val_acc: 0.076000
(Iteration 1001 / 1960) loss: 2.307522
(Iteration 1101 / 1960) loss: 2.307463
(Iteration 1201 / 1960) loss: 2.307483
(Iteration 1301 / 1960) loss: 2.307481
(Iteration 1401 / 1960) loss: 2.307486
(Epoch 3 / 4) train acc: 0.118000; val_acc: 0.078000
(Iteration 1501 / 1960) loss: 2.307491
(Iteration 1601 / 1960) loss: 2.307460
(Iteration 1701 / 1960) loss: 2.307444
(Iteration 1801 / 1960) loss: 2.307497
(Iteration 1901 / 1960) loss: 2.307519
(Epoch 4 / 4) train acc: 0.095000; val_acc: 0.082000
Iteration: lr_iteration = 2 reg_iteration = 0
(Iteration 1 / 1960) loss: 2.302625
(Epoch 0 / 4) train acc: 0.114000; val_acc: 0.108000
(Iteration 101 / 1960) loss: 2.302556
(Iteration 201 / 1960) loss: 2.302619
(Iteration 301 / 1960) loss: 2.302506
(Iteration 401 / 1960) loss: 2.302433
(Epoch 1 / 4) train acc: 0.144000; val_acc: 0.148000
(Iteration 501 / 1960) loss: 2.302529

```

```

(Iteration 601 / 1960) loss: 2.302458
(Iteration 701 / 1960) loss: 2.302504
(Iteration 801 / 1960) loss: 2.302432
(Iteration 901 / 1960) loss: 2.302394
(Epoch 2 / 4) train acc: 0.117000; val_acc: 0.105000
(Iteration 1001 / 1960) loss: 2.302247
(Iteration 1101 / 1960) loss: 2.302232
(Iteration 1201 / 1960) loss: 2.302417
(Iteration 1301 / 1960) loss: 2.302137
(Iteration 1401 / 1960) loss: 2.302031
(Epoch 3 / 4) train acc: 0.105000; val_acc: 0.102000
(Iteration 1501 / 1960) loss: 2.302074
(Iteration 1601 / 1960) loss: 2.302109
(Iteration 1701 / 1960) loss: 2.302237
(Iteration 1801 / 1960) loss: 2.302375
(Iteration 1901 / 1960) loss: 2.301977
(Epoch 4 / 4) train acc: 0.111000; val_acc: 0.093000
Iteration: lr_iteration = 2 reg_iteration = 1
(Iteration 1 / 1960) loss: 2.302594
(Epoch 0 / 4) train acc: 0.100000; val_acc: 0.109000
(Iteration 101 / 1960) loss: 2.302526
(Iteration 201 / 1960) loss: 2.302579
(Iteration 301 / 1960) loss: 2.302472
(Iteration 401 / 1960) loss: 2.302488
(Epoch 1 / 4) train acc: 0.131000; val_acc: 0.136000
(Iteration 501 / 1960) loss: 2.302465
(Iteration 601 / 1960) loss: 2.302498
(Iteration 701 / 1960) loss: 2.302430
(Iteration 801 / 1960) loss: 2.302494
(Iteration 901 / 1960) loss: 2.302364
(Epoch 2 / 4) train acc: 0.168000; val_acc: 0.152000
(Iteration 1001 / 1960) loss: 2.302391
(Iteration 1101 / 1960) loss: 2.302217
(Iteration 1201 / 1960) loss: 2.302295
(Iteration 1301 / 1960) loss: 2.302492
(Iteration 1401 / 1960) loss: 2.302423
(Epoch 3 / 4) train acc: 0.195000; val_acc: 0.190000
(Iteration 1501 / 1960) loss: 2.302113
(Iteration 1601 / 1960) loss: 2.302130
(Iteration 1701 / 1960) loss: 2.302211
(Iteration 1801 / 1960) loss: 2.302175
(Iteration 1901 / 1960) loss: 2.302161
(Epoch 4 / 4) train acc: 0.192000; val_acc: 0.178000
Iteration: lr_iteration = 2 reg_iteration = 2
(Iteration 1 / 1960) loss: 2.302670
(Epoch 0 / 4) train acc: 0.098000; val_acc: 0.085000
(Iteration 101 / 1960) loss: 2.302670
(Iteration 201 / 1960) loss: 2.302620

```

```

(Iteration 301 / 1960) loss: 2.302548
(Iteration 401 / 1960) loss: 2.302539
(Epoch 1 / 4) train acc: 0.105000; val_acc: 0.085000
(Iteration 501 / 1960) loss: 2.302585
(Iteration 601 / 1960) loss: 2.302358
(Iteration 701 / 1960) loss: 2.302467
(Iteration 801 / 1960) loss: 2.302552
(Iteration 901 / 1960) loss: 2.302427
(Epoch 2 / 4) train acc: 0.121000; val_acc: 0.099000
(Iteration 1001 / 1960) loss: 2.302409
(Iteration 1101 / 1960) loss: 2.302335
(Iteration 1201 / 1960) loss: 2.302462
(Iteration 1301 / 1960) loss: 2.302607
(Iteration 1401 / 1960) loss: 2.302406
(Epoch 3 / 4) train acc: 0.158000; val_acc: 0.113000
(Iteration 1501 / 1960) loss: 2.302603
(Iteration 1601 / 1960) loss: 2.302245
(Iteration 1701 / 1960) loss: 2.302126
(Iteration 1801 / 1960) loss: 2.302356
(Iteration 1901 / 1960) loss: 2.302412
(Epoch 4 / 4) train acc: 0.117000; val_acc: 0.084000
Iteration: lr_iteration = 2 reg_iteration = 3
(Iteration 1 / 1960) loss: 2.307516
(Epoch 0 / 4) train acc: 0.108000; val_acc: 0.115000
(Iteration 101 / 1960) loss: 2.307371
(Iteration 201 / 1960) loss: 2.307273
(Iteration 301 / 1960) loss: 2.307183
(Iteration 401 / 1960) loss: 2.307026
(Epoch 1 / 4) train acc: 0.166000; val_acc: 0.145000
(Iteration 501 / 1960) loss: 2.306863
(Iteration 601 / 1960) loss: 2.306851
(Iteration 701 / 1960) loss: 2.306660
(Iteration 801 / 1960) loss: 2.306609
(Iteration 901 / 1960) loss: 2.306467
(Epoch 2 / 4) train acc: 0.132000; val_acc: 0.129000
(Iteration 1001 / 1960) loss: 2.306415
(Iteration 1101 / 1960) loss: 2.306509
(Iteration 1201 / 1960) loss: 2.306285
(Iteration 1301 / 1960) loss: 2.306163
(Iteration 1401 / 1960) loss: 2.306170
(Epoch 3 / 4) train acc: 0.108000; val_acc: 0.083000
(Iteration 1501 / 1960) loss: 2.306041
(Iteration 1601 / 1960) loss: 2.306028
(Iteration 1701 / 1960) loss: 2.305892
(Iteration 1801 / 1960) loss: 2.305671
(Iteration 1901 / 1960) loss: 2.305934
(Epoch 4 / 4) train acc: 0.116000; val_acc: 0.078000
Iteration: lr_iteration = 3 reg_iteration = 0

```

```

(Iteration 1 / 1960) loss: 2.302578
(Epoch 0 / 4) train acc: 0.092000; val_acc: 0.102000
(Iteration 101 / 1960) loss: 2.233268
(Iteration 201 / 1960) loss: 1.735655
(Iteration 301 / 1960) loss: 1.563392
(Iteration 401 / 1960) loss: 1.373567
(Epoch 1 / 4) train acc: 0.476000; val_acc: 0.483000
(Iteration 501 / 1960) loss: 1.469317
(Iteration 601 / 1960) loss: 1.490505
(Iteration 701 / 1960) loss: 1.360365
(Iteration 801 / 1960) loss: 1.427589
(Iteration 901 / 1960) loss: 1.218399
(Epoch 2 / 4) train acc: 0.541000; val_acc: 0.516000
(Iteration 1001 / 1960) loss: 1.112765
(Iteration 1101 / 1960) loss: 1.345577
(Iteration 1201 / 1960) loss: 1.358915
(Iteration 1301 / 1960) loss: 1.410396
(Iteration 1401 / 1960) loss: 1.330399
(Epoch 3 / 4) train acc: 0.569000; val_acc: 0.527000
(Iteration 1501 / 1960) loss: 1.147417
(Iteration 1601 / 1960) loss: 1.168851
(Iteration 1701 / 1960) loss: 1.024295
(Iteration 1801 / 1960) loss: 1.268464
(Iteration 1901 / 1960) loss: 1.193003
(Epoch 4 / 4) train acc: 0.582000; val_acc: 0.562000
Iteration: lr_iteration = 3 reg_iteration = 1
(Iteration 1 / 1960) loss: 2.302576
(Epoch 0 / 4) train acc: 0.118000; val_acc: 0.106000
(Iteration 101 / 1960) loss: 2.222111
(Iteration 201 / 1960) loss: 1.783056
(Iteration 301 / 1960) loss: 1.563432
(Iteration 401 / 1960) loss: 1.555717
(Epoch 1 / 4) train acc: 0.504000; val_acc: 0.505000
(Iteration 501 / 1960) loss: 1.433862
(Iteration 601 / 1960) loss: 1.356458
(Iteration 701 / 1960) loss: 1.311791
(Iteration 801 / 1960) loss: 1.194900
(Iteration 901 / 1960) loss: 1.321054
(Epoch 2 / 4) train acc: 0.552000; val_acc: 0.510000
(Iteration 1001 / 1960) loss: 1.265947
(Iteration 1101 / 1960) loss: 1.064624
(Iteration 1201 / 1960) loss: 1.182603
(Iteration 1301 / 1960) loss: 1.283275
(Iteration 1401 / 1960) loss: 1.365324
(Epoch 3 / 4) train acc: 0.563000; val_acc: 0.534000
(Iteration 1501 / 1960) loss: 1.086229
(Iteration 1601 / 1960) loss: 1.126564
(Iteration 1701 / 1960) loss: 1.125215

```

(Iteration 1801 / 1960) loss: 1.131257  
(Iteration 1901 / 1960) loss: 1.166111  
(Epoch 4 / 4) train acc: 0.561000; val\_acc: 0.554000  
Iteration: lr\_iteration = 3 reg\_iteration = 2  
(Iteration 1 / 1960) loss: 2.302654  
(Epoch 0 / 4) train acc: 0.104000; val\_acc: 0.107000  
(Iteration 101 / 1960) loss: 2.225577  
(Iteration 201 / 1960) loss: 1.862331  
(Iteration 301 / 1960) loss: 1.473816  
(Iteration 401 / 1960) loss: 1.606194  
(Epoch 1 / 4) train acc: 0.504000; val\_acc: 0.483000  
(Iteration 501 / 1960) loss: 1.434281  
(Iteration 601 / 1960) loss: 1.323109  
(Iteration 701 / 1960) loss: 1.440046  
(Iteration 801 / 1960) loss: 1.440741  
(Iteration 901 / 1960) loss: 1.570624  
(Epoch 2 / 4) train acc: 0.515000; val\_acc: 0.512000  
(Iteration 1001 / 1960) loss: 1.307521  
(Iteration 1101 / 1960) loss: 1.220555  
(Iteration 1201 / 1960) loss: 1.482726  
(Iteration 1301 / 1960) loss: 1.352119  
(Iteration 1401 / 1960) loss: 1.198110  
(Epoch 3 / 4) train acc: 0.552000; val\_acc: 0.527000  
(Iteration 1501 / 1960) loss: 1.369568  
(Iteration 1601 / 1960) loss: 1.376530  
(Iteration 1701 / 1960) loss: 1.267542  
(Iteration 1801 / 1960) loss: 1.144453  
(Iteration 1901 / 1960) loss: 1.186016  
(Epoch 4 / 4) train acc: 0.594000; val\_acc: 0.552000  
Iteration: lr\_iteration = 3 reg\_iteration = 3  
(Iteration 1 / 1960) loss: 2.307470  
(Epoch 0 / 4) train acc: 0.109000; val\_acc: 0.098000  
(Iteration 101 / 1960) loss: 2.290977  
(Iteration 201 / 1960) loss: 2.225333  
(Iteration 301 / 1960) loss: 2.073169  
(Iteration 401 / 1960) loss: 2.059855  
(Epoch 1 / 4) train acc: 0.357000; val\_acc: 0.377000  
(Iteration 501 / 1960) loss: 1.998663  
(Iteration 601 / 1960) loss: 2.094819  
(Iteration 701 / 1960) loss: 1.955593  
(Iteration 801 / 1960) loss: 1.870003  
(Iteration 901 / 1960) loss: 1.984754  
(Epoch 2 / 4) train acc: 0.448000; val\_acc: 0.422000  
(Iteration 1001 / 1960) loss: 2.029154  
(Iteration 1101 / 1960) loss: 1.910095  
(Iteration 1201 / 1960) loss: 2.074097  
(Iteration 1301 / 1960) loss: 1.971903  
(Iteration 1401 / 1960) loss: 1.997469

```

(Epoch 3 / 4) train acc: 0.448000; val_acc: 0.435000
(Iteration 1501 / 1960) loss: 1.935737
(Iteration 1601 / 1960) loss: 1.870400
(Iteration 1701 / 1960) loss: 1.892394
(Iteration 1801 / 1960) loss: 1.935193
(Iteration 1901 / 1960) loss: 1.960940
(Epoch 4 / 4) train acc: 0.449000; val_acc: 0.413000
lr 1.000000e-07 reg 1.000000e-07 val accuracy: 0.096000
lr 1.000000e-07 reg 1.000000e-05 val accuracy: 0.090000
lr 1.000000e-07 reg 1.000000e-03 val accuracy: 0.093000
lr 1.000000e-07 reg 1.000000e-01 val accuracy: 0.103000
lr 1.000000e-05 reg 1.000000e-07 val accuracy: 0.077000
lr 1.000000e-05 reg 1.000000e-05 val accuracy: 0.090000
lr 1.000000e-05 reg 1.000000e-03 val accuracy: 0.094000
lr 1.000000e-05 reg 1.000000e-01 val accuracy: 0.082000
lr 1.000000e-03 reg 1.000000e-07 val accuracy: 0.148000
lr 1.000000e-03 reg 1.000000e-05 val accuracy: 0.190000
lr 1.000000e-03 reg 1.000000e-03 val accuracy: 0.113000
lr 1.000000e-03 reg 1.000000e-01 val accuracy: 0.145000
lr 1.000000e-01 reg 1.000000e-07 val accuracy: 0.562000
lr 1.000000e-01 reg 1.000000e-05 val accuracy: 0.554000
lr 1.000000e-01 reg 1.000000e-03 val accuracy: 0.552000
lr 1.000000e-01 reg 1.000000e-01 val accuracy: 0.435000
best validation accuracy achieved during cross-validation: 0.562000

```

```

[11]: # Run your best neural net classifier on the test set. You should be able
      # to get more than 55% accuracy.

```

```

y_test_pred = np.argmax(best_net.loss(data['X_test']), axis=1)
test_acc = (y_test_pred == data['y_test']).mean()
print(test_acc)

```

0.54